

Spring Data for Apache Solr

Christoph Strobl, Oliver Gierke, Mark Pollack, Thomas Risberg, Jay Bryant

– Version 4.0.8.RELEASE, 2019-05-13

Table of Contents

Preface

Project Metadata

Requirements

1. New & Noteworthy

1.1. What's New in Spring Data for Apache Solr 2.1

1.2. What's New in Spring Data for Apache Solr 2.0

1.3. What's New in Spring Data for Apache Solr 1.5

1.4. What's New in Spring Data for Apache Solr 1.4

2. Working with Spring Data Repositories

2.1. Core concepts

2.2. Query methods

2.3. Defining Repository Interfaces

2.3.1. Fine-tuning Repository Definition

2.3.2. Null Handling of Repository Methods

2.3.3. Using Repositories with Multiple Spring Data Modules

2.4. Defining Query Methods

2.4.1. Query Lookup Strategies

2.4.2. Query Creation

2.4.3. Property Expressions

2.4.4. Special parameter handling

2.4.5. Limiting Query Results

2.4.6. Streaming query results

2.4.7. Async query results

2.5. Creating Repository Instances

2.5.1. XML configuration

2.5.2. JavaConfig

2.5.3. Standalone usage

2.6. Custom Implementations for Spring Data Repositories

2.6.1. Customizing Individual Repositories

2.6.2. Customize the Base Repository

2.7. Publishing Events from Aggregate Roots

2.8. Spring Data Extensions

2.8.1. Querydsl Extension

2.8.2. Web support

2.8.3. Repository Populators

Reference Documentation

3. Solr Repositories

3.1. Spring Namespace

3.2. Annotation-based Configuration

3.3. Using CDI to Set up Solr Repositores

3.4. Transaction Support

3.5. Query Methods

3.5.1. Query Lookup Strategies

3.5.2. Query Creation

3.5.3. Using the `@Query` Annotation

3.5.4. Using Named Queries

3.6. Document Mapping

3.6.1. Object Mapping Fundamentals

3.6.2. `MappingSolrConverter`

4. Miscellaneous Solr Operation Support

4.1. Collection / Core Name

4.2. Partial Updates

4.3. Projection

4.4. Faceting

4.4.1. Range Faceting

4.4.2. Pivot Faceting

4.5. Terms

4.6. Result Grouping and Field Collapsing

4.7. Field Stats

4.8. Filter Query

4.9. Time Allowed for a Search

4.10. Boosting the Document Score

4.10.1. Index Time Boosts

4.11. Selecting the Request Handler

4.12. Using Joins

4.13. Highlighting

4.14. Spellchecking

4.14.1. Spellcheck Options

4.14.2. `@Spellcheck`

4.15. Using Functions

4.16. Real-time Get

4.17. Special Fields

4.17.1. @Score

4.18. Nested Documents

Appendix

Appendix A: Namespace reference

The `<repositories />` Element

Appendix B: Populators namespace reference

The `<populator />` element

Appendix C: Repository query keywords

Supported query keywords

Appendix D: Repository query return types

Supported Query Return Types

© 2012-2019 The original author(s).



Copies of this document may be made for your own use and for distribution to others, provided that you do not charge any fee for such copies and further provided that each copy contains this Copyright Notice, whether distributed in print or electronically.

Preface

The Spring Data for Apache Solr project applies core Spring concepts to the development of solutions by using the Apache Solr Search Engine. We provide a “template” as a high-level abstraction for storing and querying documents. You may notice similarities to the MongoDB support in the Spring Framework.

Project Metadata

- Version Control: <https://github.com/spring-projects/spring-data-solr>
- Bugtacker: <https://jira.spring.io/browse/DATASOLR>
- Release repository: <https://repo.spring.io/libs-release>
- Milestone repository: <https://repo.spring.io/libs-milestone>
- Snapshot repository: <https://repo.spring.io/libs-snapshot>

Requirements

Spring Data Solr requires Java 8 runtime and [Apache Solr](#) 7.7. We recommend using the latest 7.7.x version.

The following Maven dependency snippet shows how to include Apache Solr in your project:

```
<dependency>
  <groupId>org.apache.solr</groupId>
  <artifactId>solr-solrj</artifactId>
  <version>${solr.version}</version>
</dependency>
```

1. New & Noteworthy

1.1. What's New in Spring Data for Apache Solr 2.1

- Autoselect Solr core by using `SolrTemplate`.
- Assert upwards compatibility with Apache Solr 6 (including 6.3).
- Support for combined `Facet` and `Highlight` queries.
- Allow reading single-valued multi-value fields into non-collection properties.
- Use the native SolrJ schema API.

1.2. What's New in Spring Data for Apache Solr 2.0

- Upgrade to Apache Solr 5.
- Support `RequestMethod` when querying.

1.3. What's New in Spring Data for Apache Solr 1.5

- Support for [Range Faceting](#).
- Automatically prefix and suffix map keys with `@Dynamic` (See: `dynamicMappedFieldValues` in [MappingSolrConverter](#)).

1.4. What's New in Spring Data for Apache Solr 1.4

- Upgraded to recent Solr 4.10.x distribution (requires Java 7).
- Added support for [Real-time Get](#).

- Get [Field Stats](#) (max, min, sum, count, mean, missing, stddev and distinct calculations).
- Use `@Score` to automatically add projection on the document score (See: [Special Fields](#)).

2. Working with Spring Data Repositories

The goal of the Spring Data repository abstraction is to significantly reduce the amount of boilerplate code required to implement data access layers for various persistence stores.

Spring Data repository documentation and your module



This chapter explains the core concepts and interfaces of Spring Data repositories. The information in this chapter is pulled from the Spring Data Commons module. It uses the configuration and code samples for the Java Persistence API (JPA) module. You should adapt the XML namespace declaration and the types to be extended to the equivalents of the particular module that you use. “[Namespace reference](#)” covers XML configuration, which is supported across all Spring Data modules supporting the repository API. “[Repository query keywords](#)” covers the query method keywords supported by the repository abstraction in general. For detailed information on the specific features of your module, see the chapter on that module of this document.

2.1. Core concepts

The central interface in the Spring Data repository abstraction is `Repository`. It takes the domain class to manage as well as the ID type of the domain class as type arguments. This interface acts primarily as a marker interface to capture the types to work with and to help you to discover interfaces that extend this one. The `CrudRepository` provides sophisticated CRUD functionality for the entity class that is being managed.

Example 1. CrudRepository interface

```
public interface CrudRepository<T, ID extends Serializable>
    extends Repository<T, ID> {

    <S extends T> S save(S entity);           1

    Optional<T> findById(ID primaryKey);     2
```

```
    Iterable<T> findAll();           3

    long count();                   4

    void delete(T entity);          5

    boolean existsById(ID primaryKey); 6

    // ... more functionality omitted.
}
```

- 1 Saves the given entity.
- 2 Returns the entity identified by the given ID.
- 3 Returns all entities.
- 4 Returns the number of entities.
- 5 Deletes the given entity.
- 6 Indicates whether an entity with the given ID exists.



We also provide persistence technology-specific abstractions, such as `JpaRepository` or `MongoRepository`. Those interfaces extend `CrudRepository` and expose the capabilities of the underlying persistence technology in addition to the rather generic persistence technology-agnostic interfaces such as `CrudRepository`.

On top of the `CrudRepository`, there is a `PagingAndSortingRepository` abstraction that adds additional methods to ease paginated access to entities:

Example 2. PagingAndSortingRepository interface

```
public interface PagingAndSortingRepository<T, ID extends Serializable>
    extends CrudRepository<T, ID> {

    Iterable<T> findAll(Sort sort);

    Page<T> findAll(Pageable pageable);
}
```

To access the second page of `User` by a page size of 20, you could do something like the following:

```
PagingAndSortingRepository<User, Long> repository = // ... get access to a bean
Page<User> users = repository.findAll(PageRequest.of(1, 20));
```

In addition to query methods, query derivation for both count and delete queries is available. The following list shows the interface definition for a derived count query:

Example 3. Derived Count Query

```
interface UserRepository extends CrudRepository<User, Long> {

    long countByLastname(String lastname);
}
```

The following list shows the interface definition for a derived delete query:

Example 4. Derived Delete Query

```
interface UserRepository extends CrudRepository<User, Long> {

    long deleteByLastname(String lastname);

    List<User> removeByLastname(String lastname);
}
```

2.2. Query methods

Standard CRUD functionality repositories usually have queries on the underlying datastore. With Spring Data, declaring those queries becomes a four-step process:

1. Declare an interface extending `Repository` or one of its subinterfaces and type it to the domain class and ID type that it should handle, as shown in the following example:

```
interface PersonRepository extends Repository<Person, Long> { ... }
```

2. Declare query methods on the interface.

```
interface PersonRepository extends Repository<Person, Long> {

    List<Person> findByLastname(String lastname);
}
```

3. Set up Spring to create proxy instances for those interfaces, either with [JavaConfig](#) or with [XML configuration](#).

a. To use Java configuration, create a class similar to the following:

```
import org.springframework.data.jpa.repository.config.EnableJpaRepositories;

@EnableJpaRepositories
class Config {}
```

b. To use XML configuration, define a bean similar to the following:

```
<?xml version="1.0" encoding="UTF-8"?>
<beans xmlns="http://www.springframework.org/schema/beans"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xmlns:jpa="http://www.springframework.org/schema/data/jpa"
  xsi:schemaLocation="http://www.springframework.org/schema/beans
    http://www.springframework.org/schema/beans/spring-beans.xsd
    http://www.springframework.org/schema/data/jpa
    http://www.springframework.org/schema/data/jpa/spring-jpa.xsd">

  <jpa:repositories base-package="com.acme.repositories"/>

</beans>
```

The JPA namespace is used in this example. If you use the repository abstraction for any other store, you need to change this to the appropriate namespace declaration of your store module. In other words, you should exchange `jpa` in favor of, for example, `mongodb`.

+ Also, note that the JavaConfig variant does not configure a package explicitly, because the package of the annotated class is used by default. To customize the package to scan, use one of the `basePackage...` attributes of the data-store-specific repository's `@Enable${store}Repositories` -annotation.

4. Inject the repository instance and use it, as shown in the following example:

```
class SomeClient {

  private final PersonRepository repository;

  SomeClient(PersonRepository repository) {
    this.repository = repository;
  }

  void doSomething() {
    List<Person> persons = repository.findByLastname("Matthews");
  }
}
```


The sections that follow explain each step in detail:

- [Defining Repository Interfaces](#)
- [Defining Query Methods](#)
- [Creating Repository Instances](#)
- [Custom Implementations for Spring Data Repositories](#)

2.3. Defining Repository Interfaces

First, define a domain class-specific repository interface. The interface must extend `Repository` and be typed to the domain class and an ID type. If you want to expose CRUD methods for that domain type, extend `CrudRepository` instead of `Repository`.

2.3.1. Fine-tuning Repository Definition

Typically, your repository interface extends `Repository`, `CrudRepository`, or `PagingAndSortingRepository`. Alternatively, if you do not want to extend Spring Data interfaces, you can also annotate your repository interface with `@RepositoryDefinition`. Extending `CrudRepository` exposes a complete set of methods to manipulate your entities. If you prefer to be selective about the methods being exposed, copy the methods you want to expose from `CrudRepository` into your domain repository.



Doing so lets you define your own abstractions on top of the provided Spring Data Repositories functionality.

The following example shows how to selectively expose CRUD methods (`findById` and `save`, in this case):

Example 5. Selectively exposing CRUD methods

```
@NoRepositoryBean
interface MyBaseRepository<T, ID extends Serializable> extends Repository<T, ID> {

    Optional<T> findById(ID id);

    <S extends T> S save(S entity);
}

interface UserRepository extends MyBaseRepository<User, Long> {
    User findByEmailAddress(EmailAddress emailAddress);
}
```

In the prior example, you defined a common base interface for all your domain repositories and exposed `findById(...)` as well as `save(...)`. These methods are routed into the base repository implementation of the store of your choice provided by Spring Data (for example, if you use JPA, the implementation is `SimpleJpaRepository`), because they match the method signatures in `CrudRepository`. So the `UserRepository` can now save users, find individual users by ID, and trigger a query to find `Users` by email address.



The intermediate repository interface is annotated with `@NoRepositoryBean`. Make sure you add that annotation to all repository interfaces for which Spring Data should not create instances at runtime.

2.3.2. Null Handling of Repository Methods

As of Spring Data 2.0, repository CRUD methods that return an individual aggregate instance use Java 8's `Optional` to indicate the potential absence of a value. Besides that, Spring Data supports returning the following wrapper types on query methods:

- `com.google.common.base.Optional`
- `scala.Option`
- `io.vavr.control.Option`
- `javaslang.control.Option` (deprecated as `Javaslang` is deprecated)

Alternatively, query methods can choose not to use a wrapper type at all. The absence of a query result is then indicated by returning `null`. Repository methods returning collections, collection alternatives, wrappers, and streams are guaranteed never to return `null` but rather the corresponding empty representation. See “[Repository query return types](#)” for details.

Nullability Annotations

You can express nullability constraints for repository methods by using [Spring Framework's nullability annotations](#). They provide a tooling-friendly approach and opt-in `null` checks during runtime, as follows:

- `@NonNullApi`: Used on the package level to declare that the default behavior for parameters and return values is to not accept or produce `null` values.
- `@NonNull`: Used on a parameter or return value that must not be `null` (not needed on a parameter and return value where `@NonNullApi` applies).
- `@Nullable`: Used on a parameter or return value that can be `null`.

Spring annotations are meta-annotated with [JSR 305](#) annotations (a dormant but widely spread JSR). JSR 305 meta-annotations let tooling vendors such as [IDEA](#), [Eclipse](#), and [Kotlin](#) provide null-safety support in a generic way, without having to hard-code support for Spring annotations. To enable runtime checking of nullability constraints for query methods, you need to activate non-nullability on the package level by using Spring's `@NonNullApi` in `package-info.java`, as shown in the following example:

Example 6. Declaring Non-nullability in `package-info.java`

```
@org.springframework.lang.NonNullApi
package com.acme;
```

Once non-null defaulting is in place, repository query method invocations get validated at runtime for nullability constraints. If a query execution result violates the defined constraint, an exception is thrown. This happens when the method would return `null` but is declared as non-nullable (the default with the annotation defined on the package the repository resides in). If you want to opt-in to nullable results again, selectively use `@Nullable` on individual methods. Using the result wrapper types mentioned at the start of this section continues to work as expected: An empty result is translated into the value that represents absence.

The following example shows a number of the techniques just described:

Example 7. Using different nullability constraints

```
package com.acme;                                     1

import org.springframework.lang.Nullable;

interface UserRepository extends Repository<User, Long> {

    User getByEmailAddress(EmailAddress emailAddress);    2

    @Nullable
    User findByEmailAddress(@Nullable EmailAddress emailAddress);    3

    Optional<User> findOptionalByEmailAddress(EmailAddress emailAddress);    4
}
```

- 1 The repository resides in a package (or sub-package) for which we have defined non-null behavior.
- 2 Throws an `EmptyResultDataAccessException` when the query executed does not produce a result. Throws an `IllegalArgumentException` when the

- emailAddress handed to the method is null.
- 3 Returns null when the query executed does not produce a result. Also accepts null as the value for emailAddress.
- 4 Returns Optional.empty() when the query executed does not produce a result. Throws an IllegalArgumentException when the emailAddress handed to the method is null.

Nullability in Kotlin-based Repositories

Kotlin has the definition of [nullability constraints](#) baked into the language. Kotlin code compiles to bytecode, which does not express nullability constraints through method signatures but rather through compiled-in metadata. Make sure to include the `kotlin-reflect` JAR in your project to enable introspection of Kotlin's nullability constraints. Spring Data repositories use the language mechanism to define those constraints to apply the same runtime checks, as follows:

Example 8. Using nullability constraints on Kotlin repositories

```
interface UserRepository : Repository<User, String> {  
  
    fun findByUsername(username: String): User      1  
  
    fun findByFirstname(firstname: String?): User?  2  
}
```

- 1 The method defines both the parameter and the result as non-nullable (the Kotlin default). The Kotlin compiler rejects method invocations that pass null to the method. If the query execution yields an empty result, an `EmptyResultDataAccessException` is thrown.
- 2 This method accepts null for the `firstname` parameter and returns null if the query execution does not produce a result.

2.3.3. Using Repositories with Multiple Spring Data Modules

Using a unique Spring Data module in your application makes things simple, because all repository interfaces in the defined scope are bound to the Spring Data module. Sometimes, applications require using more than one Spring Data module. In such cases, a repository definition must distinguish between persistence technologies. When it detects multiple repository factories on the class path, Spring Data enters strict repository configuration mode. Strict configuration uses details on the repository or the domain class to decide about Spring Data module binding for a repository definition:

1. If the repository definition extends the module-specific repository, then it is a valid candidate for the particular Spring Data module.
2. If the domain class is annotated with the module-specific type annotation, then it is a valid candidate for the particular Spring Data module. Spring Data modules accept either third-party annotations (such as JPA's `@Entity`) or provide their own annotations (such as `@Document` for Spring Data MongoDB and Spring Data Elasticsearch).

The following example shows a repository that uses module-specific interfaces (JPA in this case):

Example 9. Repository definitions using module-specific interfaces

```
interface MyRepository extends JpaRepository<User, Long> { }

@NoRepositoryBean
interface MyBaseRepository<T, ID extends Serializable> extends JpaRepository<T, ID> {
    ...
}

interface UserRepository extends MyBaseRepository<User, Long> {
    ...
}
```

`MyRepository` and `UserRepository` extend `JpaRepository` in their type hierarchy. They are valid candidates for the Spring Data JPA module.

The following example shows a repository that uses generic interfaces:

Example 10. Repository definitions using generic interfaces

```
interface AmbiguousRepository extends Repository<User, Long> {
    ...
}

@NoRepositoryBean
interface MyBaseRepository<T, ID extends Serializable> extends CrudRepository<T, ID> {
    ...
}

interface AmbiguousUserRepository extends MyBaseRepository<User, Long> {
    ...
}
```

`AmbiguousRepository` and `AmbiguousUserRepository` extend only `Repository` and `CrudRepository` in their type hierarchy. While this is perfectly fine when using a

unique Spring Data module, multiple modules cannot distinguish to which particular Spring Data these repositories should be bound.

The following example shows a repository that uses domain classes with annotations:

Example 11. Repository definitions using domain classes with annotations

```
interface PersonRepository extends Repository<Person, Long> {  
    ...  
}  
  
@Entity  
class Person {  
    ...  
}  
  
interface UserRepository extends Repository<User, Long> {  
    ...  
}  
  
@Document  
class User {  
    ...  
}
```

PersonRepository references Person, which is annotated with the JPA @Entity annotation, so this repository clearly belongs to Spring Data JPA. UserRepository references User, which is annotated with Spring Data MongoDB's @Document annotation.

The following bad example shows a repository that uses domain classes with mixed annotations:

Example 12. Repository definitions using domain classes with mixed annotations

```
interface JpaPersonRepository extends Repository<Person, Long> {  
    ...  
}  
  
interface MongoDBPersonRepository extends Repository<Person, Long> {  
    ...  
}  
  
@Entity  
@Document  
class Person {  
    ...  
}
```

This example shows a domain class using both JPA and Spring Data MongoDB annotations. It defines two repositories, `JpaPersonRepository` and `MongoDBPersonRepository`. One is intended for JPA and the other for MongoDB usage. Spring Data is no longer able to tell the repositories apart, which leads to undefined behavior.

[Repository type details](#) and [distinguishing domain class annotations](#) are used for strict repository configuration to identify repository candidates for a particular Spring Data module. Using multiple persistence technology-specific annotations on the same domain type is possible and enables reuse of domain types across multiple persistence technologies. However, Spring Data can then no longer determine a unique module with which to bind the repository.

The last way to distinguish repositories is by scoping repository base packages. Base packages define the starting points for scanning for repository interface definitions, which implies having repository definitions located in the appropriate packages. By default, annotation-driven configuration uses the package of the configuration class. The [base package in XML-based configuration](#) is mandatory.

The following example shows annotation-driven configuration of base packages:

Example 13. Annotation-driven configuration of base packages

```
@EnableJpaRepositories(basePackages = "com.acme.repositories.jpa")
@EnableMongoRepositories(basePackages = "com.acme.repositories.mongo")
interface Configuration { }
```

2.4. Defining Query Methods

The repository proxy has two ways to derive a store-specific query from the method name:

- By deriving the query from the method name directly.
- By using a manually defined query.

Available options depend on the actual store. However, there must be a strategy that decides what actual query is created. The next section describes the available options.

2.4.1. Query Lookup Strategies

The following strategies are available for the repository infrastructure to resolve the query. With XML configuration, you can configure the strategy at the namespace through the

query-lookup-strategy attribute. For Java configuration, you can use the `queryLookupStrategy` attribute of the `Enable${store}Repositories` annotation. Some strategies may not be supported for particular datastores.

- `CREATE` attempts to construct a store-specific query from the query method name. The general approach is to remove a given set of well known prefixes from the method name and parse the rest of the method. You can read more about query construction in [“Query Creation”](#).
- `USE_DECLARED_QUERY` tries to find a declared query and throws an exception if cannot find one. The query can be defined by an annotation somewhere or declared by other means. Consult the documentation of the specific store to find available options for that store. If the repository infrastructure does not find a declared query for the method at bootstrap time, it fails.
- `CREATE_IF_NOT_FOUND` (default) combines `CREATE` and `USE_DECLARED_QUERY`. It looks up a declared query first, and, if no declared query is found, it creates a custom method name-based query. This is the default lookup strategy and, thus, is used if you do not configure anything explicitly. It allows quick query definition by method names but also custom-tuning of these queries by introducing declared queries as needed.

2.4.2. Query Creation

The query builder mechanism built into Spring Data repository infrastructure is useful for building constraining queries over entities of the repository. The mechanism strips the prefixes `find...By`, `read...By`, `query...By`, `count...By`, and `get...By` from the method and starts parsing the rest of it. The introducing clause can contain further expressions, such as a `Distinct` to set a distinct flag on the query to be created. However, the first `By` acts as delimiter to indicate the start of the actual criteria. At a very basic level, you can define conditions on entity properties and concatenate them with `And` and `Or`. The following example shows how to create a number of queries:

Example 14. Query creation from method names

```
interface PersonRepository extends Repository<User, Long> {

    List<Person> findByEmailAddressAndLastname(EmailAddress emailAddress, String
lastname);

    // Enables the distinct flag for the query
    List<Person> findDistinctPeopleByLastnameOrFirstname(String lastname, String
firstname);
    List<Person> findPeopleDistinctByLastnameOrFirstname(String lastname, String
firstname);

    // Enabling ignoring case for an individual property
    List<Person> findByLastnameIgnoreCase(String lastname);
```



```
// Enabling ignoring case for all suitable properties
List<Person> findByLastnameAndFirstnameAllIgnoreCase(String lastname, String
firstname);

// Enabling static ORDER BY for a query
List<Person> findByLastnameOrderByFirstnameAsc(String lastname);
List<Person> findByLastnameOrderByFirstnameDesc(String lastname);
}
```

The actual result of parsing the method depends on the persistence store for which you create the query. However, there are some general things to notice:

- The expressions are usually property traversals combined with operators that can be concatenated. You can combine property expressions with `AND` and `OR`. You also get support for operators such as `Between`, `LessThan`, `GreaterThan`, and `Like` for the property expressions. The supported operators can vary by datastore, so consult the appropriate part of your reference documentation.
- The method parser supports setting an `IgnoreCase` flag for individual properties (for example, `findByLastnameIgnoreCase(...)`) or for all properties of a type that supports ignoring case (usually `String` instances — for example, `findByLastnameAndFirstnameAllIgnoreCase(...)`). Whether ignoring cases is supported may vary by store, so consult the relevant sections in the reference documentation for the store-specific query method.
- You can apply static ordering by appending an `OrderBy` clause to the query method that references a property and by providing a sorting direction (`Asc` or `Desc`). To create a query method that supports dynamic sorting, see [“Special parameter handling”](#).

2.4.3. Property Expressions

Property expressions can refer only to a direct property of the managed entity, as shown in the preceding example. At query creation time, you already make sure that the parsed property is a property of the managed domain class. However, you can also define constraints by traversing nested properties. Consider the following method signature:

```
List<Person> findByAddressZipCode(ZipCode zipCode);
```

Assume a `Person` has an `Address` with a `ZipCode`. In that case, the method creates the property traversal `x.address.zipCode`. The resolution algorithm starts by interpreting the entire part (`AddressZipCode`) as the property and checks the domain class for a property with that name (uncapitalized). If the algorithm succeeds, it uses that property. If not, the algorithm splits up the source at the camel case parts from the right side into a head and a tail and tries to find the corresponding property — in our example, `AddressZip` and `Code`.

If the algorithm finds a property with that head, it takes the tail and continues building the tree down from there, splitting the tail up in the way just described. If the first split does not match, the algorithm moves the split point to the left (Address , ZipCode) and continues.

Although this should work for most cases, it is possible for the algorithm to select the wrong property. Suppose the Person class has an addressZip property as well. The algorithm would match in the first split round already, choose the wrong property, and fail (as the type of addressZip probably has no code property).

To resolve this ambiguity you can use `_` inside your method name to manually define traversal points. So our method name would be as follows:

```
List<Person> findByAddress_ZipCode(ZipCode zipCode);
```

Because we treat the underscore character as a reserved character, we strongly advise following standard Java naming conventions (that is, not using underscores in property names but using camel case instead).

2.4.4. Special parameter handling

To handle parameters in your query, define method parameters as already seen in the preceding examples. Besides that, the infrastructure recognizes certain specific types like Pageable and Sort , to apply pagination and sorting to your queries dynamically. The following example demonstrates these features:

Example 15. Using Pageable , Slice , and Sort in query methods

```
Page<User> findByLastname(String lastname, Pageable pageable);

Slice<User> findByLastname(String lastname, Pageable pageable);

List<User> findByLastname(String lastname, Sort sort);

List<User> findByLastname(String lastname, Pageable pageable);
```

The first method lets you pass an org.springframework.data.domain.Pageable instance to the query method to dynamically add paging to your statically defined query. A Page knows about the total number of elements and pages available. It does so by the infrastructure triggering a count query to calculate the overall number. As this might be expensive (depending on the store used), you can instead return a Slice . A Slice only knows about whether a next Slice is available, which might be sufficient when walking through a larger result set.

Sorting options are handled through the `Pageable` instance, too. If you only need sorting, add an `org.springframework.data.domain.Sort` parameter to your method. As you can see, returning a `List` is also possible. In this case, the additional metadata required to build the actual `Page` instance is not created (which, in turn, means that the additional count query that would have been necessary is not issued). Rather, it restricts the query to look up only the given range of entities.



To find out how many pages you get for an entire query, you have to trigger an additional count query. By default, this query is derived from the query you actually trigger.

2.4.5. Limiting Query Results

The results of query methods can be limited by using the `first` or `top` keywords, which can be used interchangeably. An optional numeric value can be appended to `top` or `first` to specify the maximum result size to be returned. If the number is left out, a result size of 1 is assumed. The following example shows how to limit the query size:

Example 16. Limiting the result size of a query with `Top` and `First`

```
User findFirstOrderByLastNameAsc();

User findTopOrderByAgeDesc();

Page<User> queryFirst10ByLastName(String lastname, Pageable pageable);

Slice<User> findTop3ByLastName(String lastname, Pageable pageable);

List<User> findFirst10ByLastName(String lastname, Sort sort);

List<User> findTop10ByLastName(String lastname, Pageable pageable);
```

The limiting expressions also support the `Distinct` keyword. Also, for the queries limiting the result set to one instance, wrapping the result into with the `Optional` keyword is supported.

If pagination or slicing is applied to a limiting query pagination (and the calculation of the number of pages available), it is applied within the limited result.

Limiting the results in combination with dynamic sorting by using a `Sort`



parameter lets you express query methods for the 'K' smallest as well as for the 'K' biggest elements.

2.4.6. Streaming query results

The results of query methods can be processed incrementally by using a Java 8 `Stream<T>` as return type. Instead of wrapping the query results in a `Stream` data store-specific methods are used to perform the streaming, as shown in the following example:

Example 17. Stream the result of a query with Java 8 `Stream<T>`

```
@Query("select u from User u")
Stream<User> findAllByCustomQueryAndStream();

Stream<User> readAllByFirstnameNotNull();

@Query("select u from User u")
Stream<User> streamAllPaged(Pageable pageable);
```



A `Stream` potentially wraps underlying data store-specific resources and must, therefore, be closed after usage. You can either manually close the `Stream` by using the `close()` method or by using a Java 7 `try-with-resources` block, as shown in the following example:

Example 18. Working with a `Stream<T>` result in a `try-with-resources` block

```
try (Stream<User> stream = repository.findAllByCustomQueryAndStream()) {
    stream.forEach(...);
}
```



Not all Spring Data modules currently support `Stream<T>` as a return type.

2.4.7. Async query results

Repository queries can be run asynchronously by using [Spring's asynchronous method execution capability](#). This means the method returns immediately upon invocation while the actual query execution occurs in a task that has been submitted to a Spring `TaskExecutor`. Asynchronous query execution is different from reactive query execution and should not be mixed. Refer to store-specific documentation for more details on reactive support. The following example shows a number of asynchronous queries:

```
@Async
Future<User> findByFirstname(String firstname); 1

@Async
CompletableFuture<User> findOneByFirstname(String firstname); 2

@Async
ListenableFuture<User> findOneByLastname(String lastname); 3
```

- 1 Use `java.util.concurrent.Future` as the return type.
- 2 Use a Java 8 `java.util.concurrent.CompletableFuture` as the return type.
- 3 Use a `org.springframework.util.concurrent.ListenableFuture` as the return type.

2.5. Creating Repository Instances

In this section, you create instances and bean definitions for the defined repository interfaces. One way to do so is by using the Spring namespace that is shipped with each Spring Data module that supports the repository mechanism, although we generally recommend using Java configuration.

2.5.1. XML configuration

Each Spring Data module includes a `repositories` element that lets you define a base package that Spring scans for you, as shown in the following example:

Example 19. Enabling Spring Data repositories via XML

```
<?xml version="1.0" encoding="UTF-8"?>
<beans:beans xmlns:beans="http://www.springframework.org/schema/beans"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xmlns="http://www.springframework.org/schema/data/jpa"
  xsi:schemaLocation="http://www.springframework.org/schema/beans
    http://www.springframework.org/schema/beans/spring-beans.xsd
    http://www.springframework.org/schema/data/jpa
    http://www.springframework.org/schema/data/jpa/spring-jpa.xsd">
```

```
<repositories base-package="com.acme.repositories" />

</beans:beans>
```

In the preceding example, Spring is instructed to scan `com.acme.repositories` and all its sub-packages for interfaces extending `Repository` or one of its sub-interfaces. For each interface found, the infrastructure registers the persistence technology-specific `FactoryBean` to create the appropriate proxies that handle invocations of the query methods. Each bean is registered under a bean name that is derived from the interface name, so an interface of `UserRepository` would be registered under `userRepository`. The `base-package` attribute allows wildcards so that you can define a pattern of scanned packages.

Using filters

By default, the infrastructure picks up every interface extending the persistence technology-specific `Repository` sub-interface located under the configured base package and creates a bean instance for it. However, you might want more fine-grained control over which interfaces have bean instances created for them. To do so, use `<include-filter />` and `<exclude-filter />` elements inside the `<repositories />` element. The semantics are exactly equivalent to the elements in Spring's context namespace. For details, see the [Spring reference documentation](#) for these elements.

For example, to exclude certain interfaces from instantiation as repository beans, you could use the following configuration:

Example 20. Using exclude-filter element

```
<repositories base-package="com.acme.repositories">
  <context:exclude-filter type="regex" expression=".*SomeRepository" />
</repositories>
```

The preceding example excludes all interfaces ending in `SomeRepository` from being instantiated.

2.5.2. JavaConfig

The repository infrastructure can also be triggered by using a store-specific `@Enable${store}Repositories` annotation on a JavaConfig class. For an introduction into Java-based configuration of the Spring container, see [JavaConfig in the Spring reference documentation](#).

A sample configuration to enable Spring Data repositories resembles the following:

Example 21. Sample annotation based repository configuration

```
@Configuration
@EnableJpaRepositories("com.acme.repositories")
class ApplicationConfiguration {

    @Bean
    EntityManagerFactory entityManagerFactory() {
        // ...
    }
}
```



The preceding example uses the JPA-specific annotation, which you would change according to the store module you actually use. The same applies to the definition of the `EntityManagerFactory` bean. See the sections covering the store-specific configuration.

2.5.3. Standalone usage

You can also use the repository infrastructure outside of a Spring container — for example, in CDI environments. You still need some Spring libraries in your classpath, but, generally, you can set up repositories programmatically as well. The Spring Data modules that provide repository support ship a persistence technology-specific `RepositoryFactory` that you can use as follows:

Example 22. Standalone usage of repository factory

```
RepositoryFactorySupport factory = ... // Instantiate factory here
UserRepository repository = factory.getRepository(UserRepository.class);
```

2.6. Custom Implementations for Spring Data Repositories

This section covers repository customization and how fragments form a composite repository.

When a query method requires a different behavior or cannot be implemented by query derivation, then it is necessary to provide a custom implementation. Spring Data repositories let you provide custom repository code and integrate it with generic CRUD abstraction and query method functionality.

2.6.1. Customizing Individual Repositories

To enrich a repository with custom functionality, you must first define a fragment interface and an implementation for the custom functionality, as shown in the following example:

Example 23. Interface for custom repository functionality

```
interface CustomizedUserRepository {  
    void someCustomMethod(User user);  
}
```

Then you can let your repository interface additionally extend from the fragment interface, as shown in the following example:

Example 24. Implementation of custom repository functionality

```
class CustomizedUserRepositoryImpl implements CustomizedUserRepository {  
  
    public void someCustomMethod(User user) {  
        // Your custom implementation  
    }  
}
```



The most important part of the class name that corresponds to the fragment interface is the `Impl` postfix.

The implementation itself does not depend on Spring Data and can be a regular Spring bean. Consequently, you can use standard dependency injection behavior to inject references to other beans (such as a `JdbcTemplate`), take part in aspects, and so on.

You can let your repository interface extend the fragment interface, as shown in the following example:

Example 25. Changes to your repository interface

```
interface UserRepository extends CrudRepository<User, Long>, CustomizedUserRepository  
{  
  
    // Declare query methods here  
}
```


Extending the fragment interface with your repository interface combines the CRUD and custom functionality and makes it available to clients.

Spring Data repositories are implemented by using fragments that form a repository composition. Fragments are the base repository, functional aspects (such as [QueryDsl](#)), and custom interfaces along with their implementation. Each time you add an interface to your repository interface, you enhance the composition by adding a fragment. The base repository and repository aspect implementations are provided by each Spring Data module.

The following example shows custom interfaces and their implementations:

Example 26. Fragments with their implementations

```
interface HumanRepository {
    void someHumanMethod(User user);
}

class HumanRepositoryImpl implements HumanRepository {

    public void someHumanMethod(User user) {
        // Your custom implementation
    }
}

interface ContactRepository {

    void someContactMethod(User user);

    User anotherContactMethod(User user);
}

class ContactRepositoryImpl implements ContactRepository {

    public void someContactMethod(User user) {
        // Your custom implementation
    }

    public User anotherContactMethod(User user) {
        // Your custom implementation
    }
}
```

The following example shows the interface for a custom repository that extends `CrudRepository`:

Example 27. Changes to your repository interface

```
interface UserRepository extends CrudRepository<User, Long>, HumanRepository,
ContactRepository {

    // Declare query methods here
}
```

Repositories may be composed of multiple custom implementations that are imported in the order of their declaration. Custom implementations have a higher priority than the base implementation and repository aspects. This ordering lets you override base repository and aspect methods and resolves ambiguity if two fragments contribute the same method signature. Repository fragments are not limited to use in a single repository interface. Multiple repositories may use a fragment interface, letting you reuse customizations across different repositories.

The following example shows a repository fragment and its implementation:

Example 28. Fragments overriding save(...)

```
interface CustomizedSave<T> {
    <S extends T> S save(S entity);
}

class CustomizedSaveImpl<T> implements CustomizedSave<T> {

    public <S extends T> S save(S entity) {
        // Your custom implementation
    }
}
```

The following example shows a repository that uses the preceding repository fragment:

Example 29. Customized repository interfaces

```
interface UserRepository extends CrudRepository<User, Long>, CustomizedSave<User> {
}

interface PersonRepository extends CrudRepository<Person, Long>,
CustomizedSave<Person> {
}
```

Configuration

If you use namespace configuration, the repository infrastructure tries to autodetect custom implementation fragments by scanning for classes below the package in which it found a repository. These classes need to follow the naming convention of appending the namespace element's `repository-impl-postfix` attribute to the fragment interface name. This postfix defaults to `Impl`. The following example shows a repository that uses the default postfix and a repository that sets a custom value for the postfix:

Example 30. Configuration example

```
<repositories base-package="com.acme.repository" />

<repositories base-package="com.acme.repository" repository-impl-postfix="MyPostfix" />
```

The first configuration in the preceding example tries to look up a class called `com.acme.repository.CustomizedUserRepositoryImpl` to act as a custom repository implementation. The second example tries to lookup `com.acme.repository.CustomizedUserRepositoryMyPostfix`.

Resolution of Ambiguity

If multiple implementations with matching class names are found in different packages, Spring Data uses the bean names to identify which one to use.

Given the following two custom implementations for the `CustomizedUserRepository` shown earlier, the first implementation is used. Its bean name is `customizedUserRepositoryImpl`, which matches that of the fragment interface (`CustomizedUserRepository`) plus the postfix `Impl`.

Example 31. Resolution of ambiguous implementations

```
package com.acme.impl.one;

class CustomizedUserRepositoryImpl implements CustomizedUserRepository {

    // Your custom implementation
}
```

```
package com.acme.impl.two;

@Component("specialCustomImpl")
class CustomizedUserRepositoryImpl implements CustomizedUserRepository {

    // Your custom implementation
}
```

If you annotate the `UserRepository` interface with `@Component("specialCustom")`, the bean name plus `Impl` then matches the one defined for the repository implementation in `com.acme.impl.two`, and it is used instead of the first one.

Manual Wiring

If your custom implementation uses annotation-based configuration and autowiring only, the preceding approach shown works well, because it is treated as any other Spring bean. If your implementation fragment bean needs special wiring, you can declare the bean and name it according to the conventions described in the [preceding section](#). The infrastructure then refers to the manually defined bean definition by name instead of creating one itself. The following example shows how to manually wire a custom implementation:

Example 32. Manual wiring of custom implementations

```
<repositories base-package="com.acme.repository" />

<beans:bean id="userRepositoryImpl" class="...">
  <!-- further configuration -->
</beans:bean>
```

2.6.2. Customize the Base Repository

The approach described in the [preceding section](#) requires customization of each repository interfaces when you want to customize the base repository behavior so that all repositories are affected. To instead change behavior for all repositories, you can create an implementation that extends the persistence technology-specific repository base class. This class then acts as a custom base class for the repository proxies, as shown in the following example:

Example 33. Custom repository base class

```
class MyRepositoryImpl<T, ID extends Serializable>
  extends SimpleJpaRepository<T, ID> {

  private final EntityManager entityManager;

  MyRepositoryImpl(JpaEntityInformation entityInformation,
                  EntityManager entityManager) {
    super(entityInformation, entityManager);

    // Keep the EntityManager around to used from the newly introduced methods.
    this.entityManager = entityManager;
  }
}
```

```

    }

    @Transactional
    public <S extends T> S save(S entity) {
        // implementation goes here
    }
}

```



The class needs to have a constructor of the super class which the store-specific repository factory implementation uses. If the repository base class has multiple constructors, override the one taking an `EntityInformation` plus a store specific infrastructure object (such as an `EntityManager` or a template class).

The final step is to make the Spring Data infrastructure aware of the customized repository base class. In Java configuration, you can do so by using the `repositoryBaseClass` attribute of the `@Enable${store}Repositories` annotation, as shown in the following example:

Example 34. Configuring a custom repository base class using JavaConfig

```

@Configuration
@EnableJpaRepositories(repositoryBaseClass = MyRepositoryImpl.class)
class ApplicationConfiguration { ... }

```

A corresponding attribute is available in the XML namespace, as shown in the following example:

Example 35. Configuring a custom repository base class using XML

```

<repositories base-package="com.acme.repository"
    base-class="...MyRepositoryImpl" />

```

2.7. Publishing Events from Aggregate Roots

Entities managed by repositories are aggregate roots. In a Domain-Driven Design application, these aggregate roots usually publish domain events. Spring Data provides an

annotation called `@DomainEvents` that you can use on a method of your aggregate root to make that publication as easy as possible, as shown in the following example:

Example 36. Exposing domain events from an aggregate root

```
class AnAggregateRoot {  
  
    @DomainEvents 1  
    Collection<Object> domainEvents() {  
        // ... return events you want to get published here  
    }  
  
    @AfterDomainEventPublication 2  
    void callbackMethod() {  
        // ... potentially clean up domain events list  
    }  
}
```

- 1 The method using `@DomainEvents` can return either a single event instance or a collection of events. It must not take any arguments.
After all events have been published, we have a method annotated with
- 2 `@AfterDomainEventPublication`. It can be used to potentially clean the list of events to be published (among other uses).

The methods are called every time one of a Spring Data repository's `save(...)` methods is called.

2.8. Spring Data Extensions

This section documents a set of Spring Data extensions that enable Spring Data usage in a variety of contexts. Currently, most of the integration is targeted towards Spring MVC.

2.8.1. Querydsl Extension

[Querydsl](#) is a framework that enables the construction of statically typed SQL-like queries through its fluent API.

Several Spring Data modules offer integration with Querydsl through `QuerydslPredicateExecutor`, as shown in the following example:

Example 37. QuerydslPredicateExecutor interface

```
public interface QuerydslPredicateExecutor<T> {  
  
    Optional<T> findById(Predicate predicate); 1  
}
```

```
Iterable<T> findAll(Predicate predicate);    2

long count(Predicate predicate);           3

boolean exists(Predicate predicate);       4

// ... more functionality omitted.
}
```

- 1 Finds and returns a single entity matching the Predicate .
- 2 Finds and returns all entities matching the Predicate .
- 3 Returns the number of entities matching the Predicate .
- 4 Returns whether an entity that matches the Predicate exists.

To make use of Querydsl support, extend `QuerydslPredicateExecutor` on your repository interface, as shown in the following example

Example 38. Querydsl integration on repositories

```
interface UserRepository extends CrudRepository<User, Long>,
QuerydslPredicateExecutor<User> {
}
```

The preceding example lets you write typesafe queries using Querydsl Predicate instances, as shown in the following example:

```
Predicate predicate = user.firstname.equalsIgnoreCase("dave")
    .and(user.lastname.startsWithIgnoreCase("mathews"));

userRepository.findAll(predicate);
```

2.8.2. Web support



This section contains the documentation for the Spring Data web support as it is implemented in the current (and later) versions of Spring Data Commons. As the newly introduced support changes many things, we kept the documentation of the former behavior in [\[web.legacy\]](#).

Spring Data modules that support the repository programming model ship with a variety of web support. The web related components require Spring MVC JARs to be on the classpath. Some of them even provide integration with [Spring HATEOAS](#). In general, the integration support is enabled by using the `@EnableSpringDataWebSupport` annotation in your JavaConfig configuration class, as shown in the following example:

Example 39. Enabling Spring Data web support

```
@Configuration
@EnableWebMvc
@EnableSpringDataWebSupport
class WebConfiguration {}
```

The `@EnableSpringDataWebSupport` annotation registers a few components we will discuss in a bit. It will also detect Spring HATEOAS on the classpath and register integration components for it as well if present.

Alternatively, if you use XML configuration, register either `SpringDataWebConfiguration` or `HateoasAwareSpringDataWebConfiguration` as Spring beans, as shown in the following example (for `SpringDataWebConfiguration`):

Example 40. Enabling Spring Data web support in XML

```
<bean class="org.springframework.data.web.config.SpringDataWebConfiguration" />

<!-- If you use Spring HATEOAS, register this one *instead* of the former -->
<bean
class="org.springframework.data.web.config.HateoasAwareSpringDataWebConfiguration" />
```

Basic Web Support

The configuration shown in the [previous section](#) registers a few basic components:

- A [DomainClassConverter](#) to let Spring MVC resolve instances of repository-managed domain classes from request parameters or path variables.
- [HandlerMethodArgumentResolver](#) implementations to let Spring MVC resolve `Pageable` and `Sort` instances from request parameters.

DomainClassConverter

The `DomainClassConverter` lets you use domain types in your Spring MVC controller method signatures directly, so that you need not manually lookup the instances through the repository, as shown in the following example:

Example 41. A Spring MVC controller using domain types in method signatures

```
@Controller
@RequestMapping("/users")
class UserController {

    @RequestMapping("/{id}")
    String showUserForm(@PathVariable("id") User user, Model model) {

        model.addAttribute("user", user);
        return "userForm";
    }
}
```

As you can see, the method receives a `User` instance directly, and no further lookup is necessary. The instance can be resolved by letting Spring MVC convert the path variable into the `id` type of the domain class first and eventually access the instance through calling `findById(...)` on the repository instance registered for the domain type.



Currently, the repository has to implement `CrudRepository` to be eligible to be discovered for conversion.

HandlerMethodArgumentResolvers for Pageable and Sort

The configuration snippet shown in the [previous section](#) also registers a `PageableHandlerMethodArgumentResolver` as well as an instance of `SortHandlerMethodArgumentResolver`. The registration enables `Pageable` and `Sort` as valid controller method arguments, as shown in the following example:

Example 42. Using Pageable as controller method argument

```
@Controller
@RequestMapping("/users")
class UserController {

    private final UserRepository repository;

    UserController(UserRepository repository) {
        this.repository = repository;
    }

    @RequestMapping
    String showUsers(Model model, Pageable pageable) {

        model.addAttribute("users", repository.findAll(pageable));
    }
}
```

```

    return "users";
  }
}

```

The preceding method signature causes Spring MVC try to derive a `Pageable` instance from the request parameters by using the following default configuration:

Table 1. Request parameters evaluated for `Pageable` instances

page	Page you want to retrieve. 0-indexed and defaults to 0.
size	Size of the page you want to retrieve. Defaults to 20.
sort	Properties that should be sorted by in the format <code>property,property(,ASC DESC)</code> . Default sort direction is ascending. Use multiple sort parameters if you want to switch directions — for example, <code>?sort=firstname&sort=lastname,asc</code> .

To customize this behavior, register a bean implementing the `PageableHandlerMethodArgumentResolverCustomizer` interface or the `SortHandlerMethodArgumentResolverCustomizer` interface, respectively. Its `customize()` method gets called, letting you change settings, as shown in the following example:

```

@Bean SortHandlerMethodArgumentResolverCustomizer sortCustomizer() {
    return s -> s.setPropertyDelimiter("<->");
}

```

If setting the properties of an existing `MethodArgumentResolver` is not sufficient for your purpose, extend either `SpringDataWebConfiguration` or the HATEOAS-enabled equivalent, override the `pageableResolver()` or `sortResolver()` methods, and import your customized configuration file instead of using the `@Enable` annotation.

If you need multiple `Pageable` or `Sort` instances to be resolved from the request (for multiple tables, for example), you can use Spring's `@Qualifier` annotation to distinguish one from another. The request parameters then have to be prefixed with `${qualifier}_`. The following example shows the resulting method signature:

```

String showUsers(Model model,
    @Qualifier("thing1") Pageable first,
    @Qualifier("thing2") Pageable second) { ... }

```

you have to populate `thing1_page` and `thing2_page` and so on.

The default `Pageable` passed into the method is equivalent to a `PageRequest.of(0, 20)` but can be customized by using the `@PageableDefault` annotation on the `Pageable` parameter.

Hypermedia Support for Pageables

Spring HATEOAS ships with a representation model class (`PagedResources`) that allows enriching the content of a `Page` instance with the necessary `Page` metadata as well as links to let the clients easily navigate the pages. The conversion of a `Page` to a `PagedResources` is done by an implementation of the Spring HATEOAS `ResourceAssembler` interface, called the `PagedResourcesAssembler`. The following example shows how to use a `PagedResourcesAssembler` as a controller method argument:

Example 43. Using a `PagedResourcesAssembler` as controller method argument

```
@Controller
class PersonController {

    @Autowired PersonRepository repository;

    @RequestMapping(value = "/persons", method = RequestMethod.GET)
    ResponseEntity<PagedResources<Person>> persons(Pageable pageable,
        PagedResourcesAssembler assembler) {

        Page<Person> persons = repository.findAll(pageable);
        return new ResponseEntity<>(assembler.toResources(persons), HttpStatus.OK);
    }
}
```

Enabling the configuration as shown in the preceding example lets the `PagedResourcesAssembler` be used as a controller method argument. Calling `toResources(...)` on it has the following effects:

- The content of the `Page` becomes the content of the `PagedResources` instance.
- The `PagedResources` object gets a `PageMetadata` instance attached, and it is populated with information from the `Page` and the underlying `PageRequest`.
- The `PagedResources` may get `prev` and `next` links attached, depending on the page's state. The links point to the URI to which the method maps. The pagination parameters added to the method match the setup of the `PageableHandlerMethodArgumentResolver` to make sure the links can be resolved later.

Assume we have 30 `Person` instances in the database. You can now trigger a request (GET <http://localhost:8080/persons>) and see output similar to the following:

```

{ "links" : [ { "rel" : "next",
                "href" : "http://localhost:8080/persons?page=1&size=20" }
            ],
  "content" : [
    ... // 20 Person instances rendered here
  ],
  "pageMetadata" : {
    "size" : 20,
    "totalElements" : 30,
    "totalPages" : 2,
    "number" : 0
  }
}

```

You see that the assembler produced the correct URI and also picked up the default configuration to resolve the parameters into a `Pageable` for an upcoming request. This means that, if you change that configuration, the links automatically adhere to the change. By default, the assembler points to the controller method it was invoked in, but that can be customized by handing in a custom `Link` to be used as base to build the pagination links, which overloads the `PagedResourcesAssembler.toResource(...)` method.

Web Databinding Support

Spring Data projections (described in [\[projections\]](#)) can be used to bind incoming request payloads by either using [JSONPath](#) expressions (requires [Jayway JsonPath](#) or [XPath](#) expressions (requires [XmlBeam](#)), as shown in the following example:

Example 44. HTTP payload binding using JSONPath or XPath expressions

```

@ProjectedPayload
public interface UserPayload {

    @XBRead("//firstname")
    @JsonPath("$.firstname")
    String getFirstname();

    @XBRead("/lastname")
    @JsonPath({ "$.lastname", "$.user.lastname" })
    String getLastName();
}

```

The type shown in the preceding example can be used as a Spring MVC handler method argument or by using `ParameterizedTypeReference` on one of `RestTemplate`'s methods. The preceding method declarations would try to find `firstname` anywhere in the given document. The `lastname` XML lookup is performed on the top-level of the incoming document. The JSON variant of that tries a top-level `lastname` first but also tries `lastname` nested in a `user` sub-document if the former does not return a value. That way, changes in

the structure of the source document can be mitigated easily without having clients calling the exposed methods (usually a drawback of class-based payload binding).

Nested projections are supported as described in [\[projections\]](#). If the method returns a complex, non-interface type, a Jackson `ObjectMapper` is used to map the final value.

For Spring MVC, the necessary converters are registered automatically as soon as `@EnableSpringDataWebSupport` is active and the required dependencies are available on the classpath. For usage with `RestTemplate`, register a `ProjectingJackson2HttpMessageConverter` (JSON) or `XmlBeamHttpMessageConverter` manually.

For more information, see the [web projection example](#) in the canonical [Spring Data Examples repository](#).

Querydsl Web Support

For those stores having [QueryDSL](#) integration, it is possible to derive queries from the attributes contained in a `Request` query string.

Consider the following query string:

```
?firstname=Dave&lastname=Matthews
```

Given the `User` object from previous examples, a query string can be resolved to the following value by using the `QuerydslPredicateArgumentResolver`.

```
QUser.user.firstname.eq("Dave").and(QUser.user.lastname.eq("Matthews"))
```



The feature is automatically enabled, along with `@EnableSpringDataWebSupport`, when `Querydsl` is found on the classpath.

Adding a `@QuerydslPredicate` to the method signature provides a ready-to-use `Predicate`, which can be run by using the `QuerydslPredicateExecutor`.



Type information is typically resolved from the method's return type. Since that information does not necessarily match the domain type, it might be a good idea to use the `root` attribute of `QuerydslPredicate`.

The following example shows how to use `@QuerydslPredicate` in a method signature:

```
@Controller
class UserController {

    @Autowired UserRepository repository;

    @RequestMapping(value = "/", method = RequestMethod.GET)
    String index(Model model, @QuerydslPredicate(root = User.class) Predicate predicate,
1         Pageable pageable, @RequestParam MultiValueMap<String, String> parameters) {

        model.addAttribute("users", repository.findAll(predicate, pageable));

        return "index";
    }
}
```

1 Resolve query string arguments to matching Predicate for User.

The default binding is as follows:

- Object on simple properties as `eq`.
- Object on collection like properties as `contains`.
- Collection on simple properties as `in`.

Those bindings can be customized through the `bindings` attribute of `@QuerydslPredicate` or by making use of Java 8 default methods and adding the `QuerydslBinderCustomizer` method to the repository interface.

```
interface UserRepository extends CrudRepository<User, String>,
    QuerydslPredicateExecutor<User>,
    QuerydslBinderCustomizer<QUser> {
1
2

    @Override
    default void customize(QuerydslBindings bindings, QUser user) {

        bindings.bind(user.username).first((path, value) -> path.contains(value))
        bindings.bind(String.class)
            .first((StringPath path, String value) -> path.containsIgnoreCase(value));
        bindings.excluding(user.password);
3
4
5
    }
}
```

- 1 QuerydslPredicateExecutor provides access to specific finder methods for Predicate.
- 2 QuerydslBinderCustomizer defined on the repository interface is automatically picked up and shortcuts @QuerydslPredicate(bindings=...) .
- 3 Define the binding for the username property to be a simple contains binding.
- 4 Define the default binding for String properties to be a case-insensitive contains match.
- 5 Exclude the password property from Predicate resolution.

2.8.3. Repository Populators

If you work with the Spring JDBC module, you are probably familiar with the support to populate a DataSource with SQL scripts. A similar abstraction is available on the repositories level, although it does not use SQL as the data definition language because it must be store-independent. Thus, the populators support XML (through Spring's OXM abstraction) and JSON (through Jackson) to define data with which to populate the repositories.

Assume you have a file `data.json` with the following content:

Example 45. Data defined in JSON

```
[ { "_class" : "com.acme.Person",
  "firstname" : "Dave",
  "lastname" : "Matthews" },
  { "_class" : "com.acme.Person",
  "firstname" : "Carter",
  "lastname" : "Beauford" } ]
```

You can populate your repositories by using the populator elements of the repository namespace provided in Spring Data Commons. To populate the preceding data to your PersonRepository, declare a populator similar to the following:

Example 46. Declaring a Jackson repository populator

```
<?xml version="1.0" encoding="UTF-8"?>
<beans xmlns="http://www.springframework.org/schema/beans"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xmlns:repository="http://www.springframework.org/schema/data/repository"
  xsi:schemaLocation="http://www.springframework.org/schema/beans
    http://www.springframework.org/schema/beans/spring-beans.xsd
    http://www.springframework.org/schema/data/repository
```

```
http://www.springframework.org/schema/data/repository/spring-repository.xsd">  
  
<repository:jackson2-populator locations="classpath:data.json" />  
  
</beans>
```

The preceding declaration causes the `data.json` file to be read and deserialized by a Jackson `ObjectMapper`.

The type to which the JSON object is unmarshalled is determined by inspecting the `_class` attribute of the JSON document. The infrastructure eventually selects the appropriate repository to handle the object that was deserialized.

To instead use XML to define the data the repositories should be populated with, you can use the `unmarshaller-populator` element. You configure it to use one of the XML marshaller options available in Spring OXM. See the [Spring reference documentation](#) for details. The following example shows how to unmarshal a repository populator with JAXB:

Example 47. Declaring an unmarshalling repository populator (using JAXB)

```
<?xml version="1.0" encoding="UTF-8"?>  
<beans xmlns="http://www.springframework.org/schema/beans"  
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"  
  xmlns:repository="http://www.springframework.org/schema/data/repository"  
  xmlns:oxm="http://www.springframework.org/schema/oxm"  
  xsi:schemaLocation="http://www.springframework.org/schema/beans  
    http://www.springframework.org/schema/beans/spring-beans.xsd  
    http://www.springframework.org/schema/data/repository  
    http://www.springframework.org/schema/data/repository/spring-repository.xsd  
    http://www.springframework.org/schema/oxm  
    http://www.springframework.org/schema/oxm/spring-oxm.xsd">  
  
  <repository:unmarshaller-populator locations="classpath:data.json"  
    unmarshaller-ref="unmarshaller" />  
  
  <oxm:jaxb2-marshaller contextPath="com.acme" />  
  
</beans>
```

Reference Documentation

3. Solr Repositories

This chapter covers details of the Solr repository implementation.

3.1. Spring Namespace

The Spring Data Solr module contains a custom namespace that allows definition of repository beans and has elements for instantiating a `SolrClient` .

Using the `repositories` element looks up Spring Data repositories as described in [Creating Repository Instances](#).

The following example shows how to set up Solr repositories that use the Spring Data Solr namespace:

Example 48. Setting up Solr repositories using the namespace

```
<?xml version="1.0" encoding="UTF-8"?>
<beans xmlns="http://www.springframework.org/schema/beans"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xmlns:solr="http://www.springframework.org/schema/data/solr"
  xsi:schemaLocation="http://www.springframework.org/schema/beans
    http://www.springframework.org/schema/beans/spring-beans.xsd
    http://www.springframework.org/schema/data/solr
    http://www.springframework.org/schema/data/solr/spring-solr.xsd">

  <solr:repositories base-package="com.acme.repositories" />
</beans>
```

Using the `solr-server` or `embedded-solr-server` element registers an instance of `SolrClient` in the context.

The following examples shows how to set up a Solr client for HTTP:

Example 49. `HttpSolrClient` using the namespace

```
<?xml version="1.0" encoding="UTF-8"?>
<beans xmlns="http://www.springframework.org/schema/beans"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xmlns:solr="http://www.springframework.org/schema/data/solr"
  xsi:schemaLocation="http://www.springframework.org/schema/beans
    http://www.springframework.org/schema/beans/spring-beans.xsd
    http://www.springframework.org/schema/data/solr
    http://www.springframework.org/schema/data/solr/spring-solr.xsd">

  <solr:solr-client id="solrClient" url="http://localhost:8983/solr" />
</beans>
```

The following example shows how to set up a load-balancing Solr client:

Example 50. LBSolrClient using the namespace

```
<?xml version="1.0" encoding="UTF-8"?>
<beans xmlns="http://www.springframework.org/schema/beans"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xmlns:solr="http://www.springframework.org/schema/data/solr"
  xsi:schemaLocation="http://www.springframework.org/schema/beans
    http://www.springframework.org/schema/beans/spring-beans.xsd
    http://www.springframework.org/schema/data/solr
    http://www.springframework.org/schema/data/solr/spring-solr.xsd">

  <solr:solr-client id="solrClient"
    url="http://localhost:8983/solr,http://localhost:8984/solr" />
</beans>
```

The following example shows how to set up an embedded Solr server:

Example 51. EmbeddedSolrServer using the namespace

```
<?xml version="1.0" encoding="UTF-8"?>
<beans xmlns="http://www.springframework.org/schema/beans"
  xmlns:xsi="http://www.w3.org/2001/XMLSchema-instance"
  xmlns:solr="http://www.springframework.org/schema/data/solr"
  xsi:schemaLocation="http://www.springframework.org/schema/beans
    http://www.springframework.org/schema/beans/spring-beans.xsd
    http://www.springframework.org/schema/data/solr
    http://www.springframework.org/schema/data/solr/spring-solr.xsd">

  <solr:embedded-solr-server id="solrClient" solrHome="classpath:com/acme/solr" />
</beans>
```

3.2. Annotation-based Configuration

The Spring Data Solr repositories support can be activated both through an XML namespace and by using an annotation through Java configuration.

The following example shows how to set up Solr repositories with Java configuration:

Example 52. Spring Data Solr repositories using Java configuration

```
@Configuration
@EnableSolrRepositories
class ApplicationConfig {

    @Bean
```

```

public SolrClient solrClient() {
    EmbeddedSolrServerFactory factory = new
EmbeddedSolrServerFactory("classpath:com/acme/solr");
    return factory.getSolrServer();
}

@Bean
public SolrOperations solrTemplate() {
    return new SolrTemplate(solrClient());
}
}

```

The preceding configuration sets up an `EmbeddedSolrServer`, which is used by the `SolrTemplate`. Spring Data Solr repositories are activated by using the `@EnableSolrRepositories` annotation, which essentially carries the same attributes as the XML namespace. If no base package is configured, it use the package in which the configuration class resides.

3.3. Using CDI to Set up Solr Repositores

You can also use CDI to set up the Spring Data Solr repositories, as the following example shows:

Example 53. Spring Data Solr repositories using Java configuration

```

class SolrTemplateProducer {

    @Produces
    @ApplicationScoped
    public SolrOperations createSolrTemplate() {
        return new SolrTemplate(new EmbeddedSolrServerFactory("classpath:com/acme/solr"));
    }
}

class ProductService {

    private ProductRepository repository;

    public Page<Product> findAvailableProductsByName(String name, Pageable pageable) {
        return repository.findByAvailableTrueAndNameStartingWith(name, pageable);
    }

    @Inject
    public void setRepository(ProductRepository repository) {
        this.repository = repository;
    }
}

```

3.4. Transaction Support

Solr support for transactions on the server level means that create, update, and delete actions since the last commit, optimization, or rollback are queued on the server and committed, optimized, or rolled back as a group. Spring Data Solr repositories participate in Spring Managed Transactions and commit or rollback changes on complete.

The following example shows how to use the `@Transactional` annotation to define a transaction (a save in this case):

```
@Transactional
public Product save(Product product) {
    Product savedProduct = jpaRepository.save(product);
    solrRepository.save(savedProduct);
    return savedProduct;
}
```

3.5. Query Methods

This section covers creating queries by using methods within Java classes.

3.5.1. Query Lookup Strategies

The Solr module supports defining a query manually as `String` or having it be derived from the method name.



There is no QueryDSL support at this time.

Declared Queries

Deriving the query from the method name is not always sufficient and may result in unreadable method names. In this case you can either use Solr named queries (see “[Using Named Queries](#)”) or use the `@Query` annotation (see “[Using the @Query Annotation](#)”).

3.5.2. Query Creation

Generally, the query creation mechanism for Solr works as described in [Query methods](#). The following example shows what a Solr query method:

Example 54. Query creation from method names

```
public interface ProductRepository extends Repository<Product, String> {
    List<Product> findByNameAndPopularity(String name, Integer popularity);
}
```

```
}

```

The preceding example translates into the following Solr query:

```
q=name:?0 AND popularity:?1

```

The following table describes the supported keywords for Solr:

Table 2. Supported keywords inside method names

Keyword	Sample	Solr Query String
And	findByNameAndPopularity	q=name:?0 AND popularity:?1
Or	findByNameOrPopularity	q=name:?0 OR popularity:?1
Is	findByName	q=name:?0
Not	findByNameNot	q=-name:?0
IsNull	findByNameIsNull	q=-name:[* TO *]
IsNotNull	findByNameIsNotNull	q=name:[* TO *]
Between	findByPopularityBetween	q=popularity:[?0 TO ?1]
LessThan	findByPopularityLessThan	q=popularity:[* TO ?0}
LessThanEqual	findByPopularityLessThanEqual	q=popularity:[* TO ?0]
GreaterThan	findByPopularityGreaterThan	q=popularity:{?0 TO *]
GreaterThanEqual	findByPopularityGreaterThanEqual	q=popularity:[?0 TO *]
Before	findByLastModifiedBefore	q=last_modified:[* TO ?0}
After	findByLastModifiedAfter	q=last_modified:{?0 TO *]
Like	findByNameLike	q=name:?0*
NotLike	findByNameNotLike	q=-name:?0*

Keyword	Sample	Solr Query String
StartingWith	findByNameStartingWith	q=name:?0*
EndingWith	findByNameEndingWith	q=name:*?0
Containing	findByNameContaining	q=name:*?0*
Matches	findByNameMatches	q=name:?0
In	findByNameIn(Collection<String> names)	q=name:(?0...)
NotIn	findByNameNotIn(Collection<String> names)	q=-name:(?0...)
Within	findByStoreWithin(Point, Distance)	q={!geofilt pt=?0.latitude,?0.longitude sfield=store d=?1}
Near	findByStoreNear(Point, Distance)	q={!bbox pt=?0.latitude,?0.longitude sfield=store d=?1}
Near	findByStoreNear(Box)	q=store[?0.start.latitude,?0.start.longitude TO ?0.end.latitude,?0.end.longitude]
True	findByAvailableTrue	q=inStock:true
False	findByAvailableFalse	q=inStock:false
OrderBy	findByAvailableTrueOrderByNameDesc	q=inStock:true&sort=name desc



Collections types can be used along with 'Like', 'NotLike', 'StartingWith', 'EndingWith' and 'Containing'.

Page<Product> findByNameLike(Collection<String> name);

3.5.3. Using the @Query Annotation

Using named queries (see “[Using Named Queries](#)”) to declare queries for entities is a valid approach and works fine for a small number of queries. As the queries themselves are tied to the Java method that executes them, you actually can bind them directly by using the Spring Data Solr `@Query` annotation. The following example uses the `@Query` annotation to declare a query:

Example 55. Declare query at the method using the `@Query` annotation.

```
public interface ProductRepository extends SolrRepository<Product, String> {  
    @Query("inStock:?0")  
    List<Product> findByAvailable(Boolean available);  
}
```

3.5.4. Using Named Queries

Named queries can be kept in a properties file and wired to the corresponding method. You should keep in mind the naming convention described in “[Query Lookup Strategies](#)” or use `@Query`. The following example shows how to declare name queries in a properties file:

Example 56. Declare named query in a properties file

```
Product.findByNamedQuery=popularity:?0  
Product.findByName=name:?0
```

The following example uses one of the named queries (`findByName`) declared in the preceding example:

```
public interface ProductRepository extends SolrCrudRepository<Product, String> {  
  
    List<Product> findByNamedQuery(Integer popularity);  
  
    @Query(name = "Product.findByName")  
    List<Product> findByAnnotatedNamedQuery(String name);  
}
```

3.6. Document Mapping

Though there is already support for Entity Mapping within SolrJ, Spring Data Solr ships with its own mapping mechanism (described in the following section).



`DocumentObjectBinder` has superior performance. Therefore, we recommend using it if you have no need for customer mapping. You can switch to `DocumentObjectBinder` by registering `SolrJConverter` within `SolrTemplate`.

3.6.1. Object Mapping Fundamentals

This section covers the fundamentals of Spring Data object mapping, object creation, field and property access, mutability and immutability. Note, that this section only applies to Spring Data modules that do not use the object mapping of the underlying data store (like JPA). Also be sure to consult the store-specific sections for store-specific object mapping, like indexes, customizing column or field names or the like.

Core responsibility of the Spring Data object mapping is to create instances of domain objects and map the store-native data structures onto those. This means we need two fundamental steps:

1. Instance creation by using one of the constructors exposed.
2. Instance population to materialize all exposed properties.

Object creation

Spring Data automatically tries to detect a persistent entity's constructor to be used to materialize objects of that type. The resolution algorithm works as follows:

1. If there's a no-argument constructor, it will be used. Other constructors will be ignored.
2. If there's a single constructor taking arguments, it will be used.
3. If there are multiple constructors taking arguments, the one to be used by Spring Data will have to be annotated with `@PersistenceConstructor`.

The value resolution assumes constructor argument names to match the property names of the entity, i.e. the resolution will be performed as if the property was to be populated, including all customizations in mapping (different datastore column or field name etc.). This also requires either parameter names information available in the class file or an `@ConstructorProperties` annotation being present on the constructor.

The value resolution can be customized by using Spring Framework's `@Value` value annotation using a store-specific SpEL expression. Please consult the section on store specific mappings for further details.

Object creation internals

To avoid the overhead of reflection, Spring Data object creation uses a factory class generated at runtime by default, which will call the domain classes constructor directly. I.e. for this example type:

```
class Person {  
    Person(String firstname, String lastname) { ... }  
}
```

we will create a factory class semantically equivalent to this one at runtime:

```
class PersonObjectInstantiator implements ObjectInstantiator {  
  
    Object newInstance(Object... args) {  
        return new Person((String) args[0], (String) args[1]);  
    }  
}
```

This gives us a roundabout 10% performance boost over reflection. For the domain class to be eligible for such optimization, it needs to adhere to a set of constraints:

- it must not be a private class
- it must not be a non-static inner class
- it must not be a CGLib proxy class
- the constructor to be used by Spring Data must not be private

If any of these criteria match, Spring Data will fall back to entity instantiation via reflection.

Property population

Once an instance of the entity has been created, Spring Data populates all remaining persistent properties of that class. Unless already populated by the entity's constructor (i.e. consumed through its constructor argument list), the identifier property will be populated first to allow the resolution of cyclic object references. After that, all non-transient properties that have not already been populated by the constructor are set on the entity instance. For that we use the following algorithm:

1. If the property is immutable but exposes a wither method (see below), we use the wither to create a new entity instance with the new property value.
2. If property access (i.e. access through getters and setters) is defined, we're invoking the setter method.
3. By default, we set the field value directly.

Property population internals

Similarly to our [optimizations in object construction](#) we also use Spring Data runtime generated accessor classes to interact with the entity instance.

```
class Person {

    private final Long id;
    private String firstname;
    private @AccessType(Type.PROPERTY) String lastname;

    Person() {
        this.id = null;
    }

    Person(Long id, String firstname, String lastname) {
        // Field assignments
    }

    Person withId(Long id) {
        return new Person(id, this.firstname, this.lastname);
    }

    void setLastname(String lastname) {
        this.lastname = lastname;
    }
}
```

Example 57. A generated Property Accessor

```
class PersonPropertyAccessor implements PersistentPropertyAccessor {

    private static final MethodHandle firstname;           2

    private Person person;                                1

    public void setProperty(PersistentProperty property, Object value) {

        String name = property.getName();

        if ("firstname".equals(name)) {
            firstname.invoke(person, (String) value);      2
        } else if ("id".equals(name)) {
```

```

    this.person = person.withId((Long) value);           3
  } else if ("lastname".equals(name)) {
    this.person.setLastname((String) value);           4
  }
}
}

```

1 PropertyAccessor's hold a mutable instance of the underlying object. This is, to enable mutations of otherwise immutable properties.

2 By default, Spring Data uses field-access to read and write property values. As per visibility rules of private fields, MethodHandles are used to interact with fields.

3 The class exposes a withId(...) method that's used to set the identifier, e.g. when an instance is inserted into the datastore and an identifier has been generated. Calling withId(...) creates a new Person object. All subsequent mutations will take place in the new instance leaving the previous untouched.

4 Using property-access allows direct method invocations without using MethodHandles .

This gives us a roundabout 25% performance boost over reflection. For the domain class to be eligible for such optimization, it needs to adhere to a set of constraints:

- Types must not reside in the default or under the java package.
- Types and their constructors must be public
- Types that are inner classes must be static .
- The used Java Runtime must allow for declaring classes in the originating ClassLoader . Java 9 and newer impose certain limitations.

By default, Spring Data attempts to use generated property accessors and falls back to reflection-based ones if a limitation is detected.

Let's have a look at the following entity:

Example 58. A sample entity

```

class Person {
    private final @Id Long id;           1
    private final String firstname, lastname; 2
}

```

```

private final LocalDate birthday;
private final int age; 3

private String comment; 4
private @AccessType(Type.PROPERTY) String remarks; 5

static Person of(String firstname, String lastname, LocalDate birthday) { 6

    return new Person(null, firstname, lastname, birthday,
        Period.between(birthday, LocalDate.now()).getYears());
}

Person(Long id, String firstname, String lastname, LocalDate birthday, int age) {
6

    this.id = id;
    this.firstname = firstname;
    this.lastname = lastname;
    this.birthday = birthday;
    this.age = age;
}

Person withId(Long id) { 1
    return new Person(id, this.firstname, this.lastname, this.birthday);
}

void setRemarks(String remarks) { 5
    this.remarks = remarks;
}
}

```

The identifier property is final but set to `null` in the constructor. The class exposes a `withId(...)` method that's used to set the identifier, e.g. when an instance is inserted into the datastore and an identifier has been generated. The original `Person` instance stays unchanged as a new one is created. The same pattern is usually applied for other properties that are store managed but might have to be changed for persistence operations.

The `firstname` and `lastname` properties are ordinary immutable properties potentially exposed through getters.

The `age` property is an immutable but derived one from the `birthday` property. With the design shown, the database value will trump the defaulting as Spring Data uses the only declared constructor. Even if the intent is that the calculation should be preferred, it's important that this constructor also takes `age` as parameter (to potentially ignore it) as otherwise the property population step will attempt to set the `age` field and fail due to it being immutable and no wither being present.

The `comment` property is mutable is populated by setting its field directly.

The `remarks` properties are mutable and populated by setting the `comment` field

directly or by invoking the setter method for

The class exposes a factory method and a constructor for object creation. The core idea here is to use factory methods instead of additional constructors to avoid the need for constructor disambiguation through `@PersistenceConstructor`. Instead, defaulting of properties is handled within the factory method.

General recommendations

- *Try to stick to immutable objects* — Immutable objects are straightforward to create as materializing an object is then a matter of calling its constructor only. Also, this avoids your domain objects to be littered with setter methods that allow client code to manipulate the objects state. If you need those, prefer to make them package protected so that they can only be invoked by a limited amount of co-located types. Constructor-only materialization is up to 30% faster than properties population.
- *Provide an all-args constructor* — Even if you cannot or don't want to model your entities as immutable values, there's still value in providing a constructor that takes all properties of the entity as arguments, including the mutable ones, as this allows the object mapping to skip the property population for optimal performance.
- *Use factory methods instead of overloaded constructors to avoid `@PersistenceConstructor`* — With an all-argument constructor needed for optimal performance, we usually want to expose more application use case specific constructors that omit things like auto-generated identifiers etc. It's an established pattern to rather use static factory methods to expose these variants of the all-args constructor.
- *Make sure you adhere to the constraints that allow the generated instantiator and property accessor classes to be used* —
- *For identifiers to be generated, still use a final field in combination with a wither method* —
- *Use Lombok to avoid boilerplate code* — As persistence operations usually require a constructor taking all arguments, their declaration becomes a tedious repetition of boilerplate parameter to field assignments that can best be avoided by using Lombok's `@AllArgsConstructor`.

Kotlin support

Spring Data adapts specifics of Kotlin to allow object creation and mutation.

Kotlin object creation

Kotlin classes are supported to be instantiated, all classes are immutable by default and require explicit property declarations to define mutable properties. Consider the following data class `Person`:

```
data class Person(val id: String, val name: String)
```

The class above compiles to a typical class with an explicit constructor. We can customize this class by adding another constructor and annotate it with `@PersistenceConstructor` to indicate a constructor preference:

```
data class Person(var id: String, val name: String) {  
    @PersistenceConstructor  
    constructor(id: String) : this(id, "unknown")  
}
```

Kotlin supports parameter optionality by allowing default values to be used if a parameter is not provided. When Spring Data detects a constructor with parameter defaulting, then it leaves these parameters absent if the data store does not provide a value (or simply returns `null`) so Kotlin can apply parameter defaulting. Consider the following class that applies parameter defaulting for `name`

```
data class Person(var id: String, val name: String = "unknown")
```

Every time the `name` parameter is either not part of the result or its value is `null`, then the name defaults to `unknown`.

Property population of Kotlin data classes

In Kotlin, all classes are immutable by default and require explicit property declarations to define mutable properties. Consider the following data class `Person`:

```
data class Person(val id: String, val name: String)
```

This class is effectively immutable. It allows to create new instances as Kotlin generates a `copy(...)` method that creates new object instances copying all property values from the existing object and applying property values provided as arguments to the method.

3.6.2. MappingSolrConverter

MappingSolrConverter lets you register custom converters for your SolrDocument and SolrInputDocument as well as for other types nested within your beans. The converter is not 100% compatible with DocumentObjectBinder, and @Indexed has to be added with readonly=true to ignore fields from being written to Solr. The following example maps a number of fields within a document:

Example 59. Sample Document Mapping

```
public class Product {
    @Field
    private String simpleProperty;

    @Field("somePropertyName")
    private String namedProperty;

    @Field
    private List<String> listOfValues;

    @Indexed(readonly = true)
    @Field("property_*")
    private List<String> ignoredFromWriting;

    @Field("mappedField_*")
    private Map<String, List<String>> mappedFieldValues;

    @Dynamic
    @Field("dynamicMappedField_*")
    private Map<String, String> dynamicMappedFieldValues;

    @Field
    private GeoLocation location;
}
```

The following table describes the properties you can map with MappingSolrConverter :

Property	Write Mapping
simpleProperty	<field name="simpleProperty">value</field>
namedProperty	<field name="somePropertyName">value</field>
listOfValues	<field name="listOfValues">value 1</field> <field name="listOfValues">value 2</field> <field name="listOfValues">value 3</field>
ignoredFromWriting	//not written to document

Property	Write Mapping
mappedFieldValues	<pre><field name="mapentry[0].key">mapentry[0].value[0] </field> <field name="mapentry[0].key">mapentry[0].value[1]</field> <field name="mapentry[1].key">mapentry[1].value[0] </field></pre>
dynamicMappedFieldValues	<pre><field name="'dynamicMappedField_' + mapentry[0].key">mapentry[0].value[0]</field> <field name="'dynamicMappedField_' + mapentry[0].key">mapentry[0].value[1]</field> <field name="'dynamicMappedField_' + mapentry[1].key">mapentry[1].value[0]</field></pre>
location	<pre><field name="location">48.362893,14.534437</field></pre>

You can register a custom converter by adding `CustomConversions` to `SolrTemplate` and initializing it with your own `Converter` implementation, as the following example shows:

```
<bean id="solrConverter"
class="org.springframework.data.solr.core.convert.MappingSolrConverter">
  <constructor-arg>
    <bean
class="org.springframework.data.solr.core.mapping.SimpleSolrMappingContext" />
  </constructor-arg>
  <property name="customConversions" ref="customConversions" />
</bean>

<bean id="customConversions"
class="org.springframework.data.solr.core.convert.SolrCustomConversions">
  <constructor-arg>
    <list>
      <bean class="com.acme.MyBeanToSolrInputDocumentConverter" />
    </list>
  </constructor-arg>
</bean>

<bean id="solrTemplate" class="org.springframework.data.solr.core.SolrTemplate">
  <constructor-arg ref="solrClient" />
  <property name="solrConverter" ref="solrConverter" />
</bean>
```


4. Miscellaneous Solr Operation Support

This chapter covers additional support for Solr operations (such as faceting) that cannot be directly accessed via the repository interface. It is recommended to add those operations as custom implementation as described in [Custom Implementations for Spring Data Repositories](#).

4.1. Collection / Core Name

Using the `@SolrDocument` annotation it is possible to customize the used collection name by either giving it a static value or use [SpEL](#) for dynamic evaluation.

```
@SolrDocument(collection = "techproducts")
class StaticCollectionName { ... }

@SolrDocument(collection = "#{@someBean.getCollectionName()}")
class DynamicCollectionName { ... }
```



The the type annotated with `@SolrDocument` is available via the `targetType` variable within the expression.

4.2. Partial Updates

PartialUpdates can be done using `PartialUpdate` which implements `Update`.

```
PartialUpdate update = new PartialUpdate("id", "123");
update.add("name", "updated-name");
solrTemplate.saveBean("collection-1", update);
```

4.3. Projection

Projections can be applied via `@Query` using the `fields` value.

```
@Query(fields = { "name", "id" })
List<ProductBean> findByNameStartingWith(String name);
```

4.4. Faceting

Faceting cannot be directly applied by using the `SolrRepository`, but the `SolrTemplate` has support for this feature. The following example shows a facet query:

```
FacetQuery query = new SimpleFacetQuery(new
Criteria(Criteria.WILDCARD).expression(Criteria.WILDCARD))
    .setFacetOptions(new FacetOptions().addFacetOnField("name").setFacetLimit(5));
FacetPage<Product> page = solrTemplate.queryForFacetPage("collection-1", query,
Product.class);
```

Facets on fields or queries can also be defined by using `@Facet`. Keep in mind that the result is a `FacetPage`.



Using `@Facet` lets you define place holders that use your input parameter as a value.

The following example uses the `@Facet` annotation to define a facet query:

```
@Query(value = ".*.*")
@Facet(fields = { "name" }, limit = 5)
FacetPage<Product> findAllFacetOnName(Pageable page);
```

The following example shows another facet query, with a prefix:

```
@Query(value = "popularity:?0")
@Facet(fields = { "name" }, limit = 5, prefix="?1")
FacetPage<Product> findByPopularityFacetOnName(int popularity, String prefix, Pageable
page);
```

Solr allows definition of facet parameters on a per field basis. In order to add special facet options to defined fields, use `FieldWithFacetParameters`, as the following example shows:

```
// produces: f.name.facet.prefix=spring
FacetOptions options = new FacetOptions();
```

```
options.addFacetOnField(new FieldWithFacetParameters("name").setPrefix("spring"));
```

4.4.1. Range Faceting

You can create range faceting queries by configuring required ranges on `FacetOptions`. You can request ranges by creating a `FacetOptions` instance, setting the options to a `FacetQuery`, and querying for a facet page through `SolrTemplate`, as follows.

```
FacetOptions facetOptions = new FacetOptions()
    .addFacetByRange(
        new FieldWithNumericRangeParameters("price", 5, 20, 5)
            .setHardEnd(true)
            .setInclude(FacetRangeInclude.ALL)
    )
    .addFacetByRange(
        new FieldWithDateRangeParameters("release", new Date(1420070400), new
Date(946684800), "+1YEAR")
            .setInclude(FacetRangeInclude.ALL)
            .setOther(FacetRangeOther.BEFORE)
    );
facetOptions.setFacetMinCount(0);

Criteria criteria = new SimpleStringCriteria("*:");
SimpleFacetQuery facetQuery = new
SimpleFacetQuery(criteria).setFacetOptions(facetOptions);
FacetPage<ExampleSolrBean> statResultPage =
solrTemplate.queryForFacetPage("collection-1", facetQuery, ExampleSolrBean.class);
```

There are two implementations of fields for facet range requests:

- **Numeric Facet Range:** Used to perform range faceting over numeric fields. To request range faceting, you can use an instance of the `org.springframework.data.solr.core.query.FacetOptions.FieldWithNumericRangeParameters` class. Its instantiation requires a field name, a start value (number), an end value (number), and a gap (number);
- **Date Facet Range:** Used to perform range faceting over date fields. To request range faceting, you can use an instance of the `org.springframework.data.solr.core.query.FacetOptions.FieldWithDateRangeParameters` class. Its instantiation requires a field name, a start value (date), an end value (date), and a gap (string). You can define the gap for this kind of field by using `org.apache.solr.util.DateMathParser` (for example, `+6MONTHS+3DAYS/DAY` means six months and three days in the future, rounded down to the nearest day).

Additionally, the following properties can be configured for a field with range parameters (`org.springframework.data.solr.core.query.FacetOptions.FieldWithRangeParameters`):

- **Hard End:** `setHardEnd(Boolean)` defines whether the last range should be abruptly ended even if the end does not satisfy $(\text{start} - \text{end}) \% \text{gap} = 0$.
- **Include:** `setInclude(org.apache.solr.common.params.FacetParams.FacetRangeInclude)` defines how boundaries (lower and upper) should be handled (exclusive or inclusive) on range facet requests.
- **Other:** `setOther(org.apache.solr.common.params.FacetParams.FacetRangeOther)` defines the additional (other) counts for the range facet (such as count of documents that are before the start of the range facet, after the end of the range facet, or even between the start and the end).

4.4.2. Pivot Faceting

Pivot faceting (decision tree) is also supported and can be queried by using `@Facet` annotation, as follows:

```
public interface {

    @Facet(pivots = @Pivot({ "category", "dimension" }, pivotMinCount = 0))
    FacetPage<Product> findByTitle(String title, Pageable page);

    @Facet(pivots = @Pivot({ "category", "dimension" }))
    FacetPage<Product> findByDescription(String description, Pageable page);

}
```

Alternatively, it can be queried by using `SolrTemplate`, as follows:

```
FacetQuery facetQuery = new SimpleFacetQuery(new SimpleStringCriteria("title:foo"));
FacetOptions facetOptions = new FacetOptions();
facetOptions.setFacetMinCount(0);
facetOptions.addFacetOnPivot("category", "dimension");
facetQuery.setFacetOptions(facetOptions);
FacetPage<Product> facetResult = solrTemplate.queryForFacetPage("collection-1",
    facetQuery, Product.class);
```

In order to retrieve the pivot results, use the `getPivot` method, as follows:

```
List<FacetPivotFieldEntry> pivot = facetResult.getPivot(new  
SimplePivotField("categories","available"));
```

4.5. Terms

A terms vector cannot directly be used within `SolrRepository` but can be applied through `SolrTemplate`. Keep in mind that the result is a `TermsPage`. The following example shows how to create a terms query:

```
TermsQuery query = SimpleTermsQuery.queryBuilder().fields("name").build();  
TermsPage page = solrTemplate.queryForTermsPage("collection-1", query);
```

4.6. Result Grouping and Field Collapsing

Result grouping cannot directly be used within `SolrRepository` but can be applied through `SolrTemplate`. Keep in mind that the result is a `GroupPage`. The following example shows how to create a result group:

```
Field field = new SimpleField("popularity");  
Function func = ExistsFunction.exists("description");  
Query query = new SimpleQuery("inStock:true");  
  
SimpleQuery groupQuery = new SimpleQuery(new SimpleStringCriteria("*:*"));  
GroupOptions groupOptions = new GroupOptions()  
    .addGroupByField(field)  
    .addGroupByFunction(func)  
    .addGroupByQuery(query);  
groupQuery.setGroupOptions(groupOptions);  
  
GroupPage<Product> page = solrTemplate.queryForGroupPage("collection-1", query,  
Product.class);  
  
GroupResult<Product> fieldGroup = page.getGroupResult(field);  
GroupResult<Product> funcGroup = page.getGroupResult(func);  
GroupResult<Product> queryGroup = page.getGroupResult(query);
```

4.7. Field Stats

Field stats are used to retrieve statistics (max, min, sum, count, mean, missing, stddev, and distinct calculations) of given fields from Solr. You can provide `StatsOptions` to

your query and read the `FieldStatsResult` from the returned `StatsPage`. You could do so, for instance, by using `SolrTemplate`, as follows:

```
// simple field stats
StatsOptions statsOptions = new StatsOptions().addField("price");

// query
SimpleQuery statsQuery = new SimpleQuery("*:");
statsQuery.setStatsOptions(statsOptions);
StatsPage<Product> statsPage = solrTemplate.queryForStatsPage("collection-1",
statsQuery, Product.class);

// retrieving stats info
FieldStatsResult priceStatResult = statResultPage.getFieldStatsResult("price");
Object max = priceStatResult.getMax();
Long missing = priceStatResult.getMissing();
```

You could achieve the same result by annotating the repository method with `@Stats`, as follows:

```
@Query("name:?0")
@Stats(value = { "price" })
StatsPage<Product> findByName(String name, Pageable page);
```

Distinct calculation and faceting are also supported:

```
// for distinct calculation
StatsOptions statsOptions = new StatsOptions()
    .addField("category")
    // for distinct calculation
    .setCalcDistinct(true)
    // for faceting
    .addFacet("availability");

// query
SimpleQuery statsQuery = new SimpleQuery("*:");
statsQuery.setStatsOptions(statsOptions);
StatsPage<Product> statsPage = solrTemplate.queryForStatsPage("collection-1",
statsQuery, Product.class);

// field stats
FieldStatsResult categoryStatResult = statResultPage.getFieldStatsResult("category");

// retrieving distinct
List<Object> categoryValues = priceStatResult.getDistinctValues();
```

```

Long distinctCount = categoryStatResult.getDistinctCount();

// retrieving faceting
Map<String, StatsResult> availabilityFacetResult =
categoryStatResult.getFacetStatsResult("availability");
Long availableCount = availabilityFacetResult.get("true").getCount();

```

The annotated (and consequently much shorter) version of the preceding example follows:

```

@Query("name:?0")
@Stats(value = "category", facets = { "availability" }, calcDistinct = true)
StatsPage<Product> findByName(String name);

```

In order to perform a selective faceting or selective distinct calculation, you can use `@SelectiveStats`, as follows:

```

// selective distinct faceting
...
Field facetField = getFacetField();
StatsOptions statsOptions = new StatsOptions()
    .addField("price")
    .addField("category").addSelectiveFacet("name").addSelectiveFacet(facetField);
...
// or annotating repository method as follows
...
@Stats(value = "price", selective = @SelectiveStats(field = "category", facets = {
    "name", "available" }))
...

// selective distinct calculation
...
StatsOptions statsOptions = new StatsOptions()
    .addField("price")
    .addField("category").setSelectiveCalcDistinct(true);
...
// or annotating repository method as follows
...
@Stats(value = "price", selective = @SelectiveStats(field = "category", calcDistinct =
true))
...

```

4.8. Filter Query

Filter Queries improve query speed and do not influence the document score. We recommend implementing geospatial search as a filter query.



In Solr, unless otherwise specified, all units of distance are kilometers and points are in degrees of latitude and longitude.

The following example shows a filter query for a geographical point (in Austria, in this case):

```
Query query = new SimpleQuery(new
Criteria("category").is("supercalifragilisticexpialidocious"));
FilterQuery fq = new SimpleFilterQuery(new Criteria("store")
.near(new Point(48.305478, 14.286699), new Distance(5)));
query.addFilterQuery(fq);
```

You can also define simple filter queries by using `@Query`.



Using `@Query` lets you define place holders that use your input parameter as a value.

The following example shows a query with placeholders (`:`):

```
@Query(value = "*:.*", filters = { "inStock:true", "popularity:[* TO 3]" })
List<Product> findAllFilterAvailableTrueAndPopularityLessThanEqual3();
```

4.9. Time Allowed for a Search

You can set the time allowed for a search to finish. This value only applies to the search and not to requests in general. Time is in milliseconds. Values less than or equal to zero imply no time restriction. Partial results may be returned, if there are any. The following example restricts the time for a search to 100 milliseconds:

```
Query query = new SimpleQuery(new SimpleStringCriteria("field_1:value_1"));
// Allowing maximum of 100ms for this search
query.setTimeAllowed(100);
```


4.10. Boosting the Document Score

You can boost the document score for matching criteria to influence the result order. You can do so either by setting `boost` on `Criteria` or by using `@Boost` for derived queries. The following example boosts the `name` parameter of the `findByNameOrDescription` query:

```
Page<Product> findByNameOrDescription(@Boost(2) String name, String description);
```

4.10.1. Index Time Boosts

Both document-based and field-based index time boosting have been removed from Apache Solr 7 and, therefore, from Spring Data for Apache Solr 4.x.

4.11. Selecting the Request Handler

You can select the request handler through the `qt` Parameter directly in `Query` or by adding `@Query` to your method signature. The following example does so by adding `@Query`:

```
@Query(requestHandler = "/instock")
Page<Product> findByNameOrDescription(String name, String description);
```

4.12. Using Joins

You can use joins within one Solr core by defining a `Join` attribute of a `Query`.



Join is not available prior to Solr 4.x.

The following example shows how to use a join:

```
SimpleQuery query = new SimpleQuery(new SimpleStringCriteria("text:ipod"));
query.setJoin(Join.from("manu_id_s").to("id"));
```

4.13. Highlighting

To highlight matches in search result, you can add `HighlightOptions` to the `SimpleHighlightQuery`. Providing `HighlightOptions` without any further attributes applies highlighting on all fields within a `SolrDocument`.



You can set field-specific highlight parameters by adding `FieldWithHighlightParameters` to `HighlightOptions`.

The following example sets highlighting for all fields in the query:

```
SimpleHighlightQuery query = new SimpleHighlightQuery(new
SimpleStringCriteria("name:with"));
query.setHighlightOptions(new HighlightOptions());
HighlightPage<Product> page = solrTemplate.queryForHighlightPage("collection-1",
query, Product.class);
```

Not all parameters are available through setters and getters but can be added directly.

The following example sets highlighting on two fields:

```
SimpleHighlightQuery query = new SimpleHighlightQuery(new
SimpleStringCriteria("name:with"));
query.setHighlightOptions(new
HighlightOptions().addHighlightParameter("hl.bs.country", "at"));
```

To apply Highlighting to derived queries, you can use `@Highlight`. If no fields are defined, highlighting is applied on all fields.

```
@Highlight(prefix = "<b>", postfix = "</b>")
HighlightPage<Product> findByName(String name, Pageable page);
```

4.14. Spellchecking

Spellchecking offers search term suggestions based on the actual query. See the [Solr Reference](#) for more details.

4.14.1. Spellcheck Options

Spellcheck query parameters are added to a request when `SpellcheckOptions` has been set, as the following example shows:

```
SimpleQuery q = new SimpleQuery("name:gren");
q.setSpellcheckOptions(SpellcheckOptions.spellcheck()           1
    .dictionaries("dict1", "dict2")                             2
    .count(5)                                                    3
    .extendedResults());                                         4
q.setRequestHandler("/spell");                                   5

SpellcheckedPage<Product> found = template.query(q, Product.class); 6
```

- 1 Enable spellcheck by setting `SpellcheckOptions`. Sets the `spellcheck=on` request parameter.
- 2 Set up the dictionaries to use for lookup.
- 3 Set the maximum number of suggestions to return.
- 4 Enable extended results, including term frequency and others.
- 5 Set the request handler, which must be capable of processing suggestions.
- 6 Run the query.

4.14.2. @Spellcheck

The `@Spellcheck` annotation allows usage of the spellcheck feature on `Repository` level. The following example shows how to use it:

```
public interface ProductRepository extends Repository<Product, String> {

    @Query(requestHandler = "/spell")
    @Spellcheck(dictionaries = { "dict1", "dic2" }, count=5, extendedResults = true)
    SpellcheckedPage<Product> findByName(String name, Pageable page);

}
```

4.15. Using Functions

Solr supports several functional expressions within queries and includes a number of functions. You can add custom functions by implementing `Function`. The following table lists which functions are supported:

Table 3. Functions

Class	Solr Function
CurrencyFunction	currency(field_name,[CODE])
DefaultValueFunction	def(field function,defaultValue)
DistanceFunction	dist(power, pointA, pointB)
DivideFunction	div(x,y)
ExistsFunction	exists(field function)
GeoDistanceFunction	geodist(sfield, latitude, longitude)
GeoHashFunction	geohash(latitude, longitude)
IfFunction	if(value field function,trueValue,falseValue)
MaxFunction	max(field function,value)
NotFunction	not(field function)
ProductFunction	product(x,y,...)
QueryFunction	query(x)
TermFrequencyFunction	termfreq(field,term)

The following example uses a QueryFunction:

```
SimpleQuery query = new SimpleQuery(new SimpleStringCriteria("text:ipod"));
query.addFilterQuery(new
    FilterQuery(Criteria.where(QueryFunction.query("name:sol*"))));
```

4.16. Real-time Get

Real-time get allows retrieval of the latest version of any document by using a unique key, without the need to reopen searchers.



Real-time get relies on the update log feature.

The following example shows a real-time get:

Example 60. Real-time get

```
Optional<Product> product = solrTemplate.getById("collection-1", "123",  
Product.class);
```

You can retrieve multiple documents by providing a collection of `ids`, as follows:

Example 61. Realtime multi-get

```
Collection<String> ids = Arrays.asList("123", "134");  
Collection<Product> products = solrTemplate.getByIds("collection-1", ids,  
Product.class);
```

4.17. Special Fields

Solr includes a number of special fields, including a score field.

4.17.1. @Score

In order to load score information of a query result, you can add a field annotated with the `@Score` annotation, indicating the property holds the document's score.



The score property needs to be numerical and can only appear once per document.

The following example shows a document with a score field:

```
public class MyEntity {  
  
    @Id  
    private String id;  
  
    @Score  
    private Float score;  
  
    // setters and getters ...  
  
}
```

4.18. Nested Documents

Nested documents allow for documents inside of other documents in a parent-child relationship.

The nested documents need to be indexed along with the parent one and cannot be updated individually. However, nested documents appear as individual documents in the index. Resolving the parent-child relation is done at query time.

To indicate that a property should be treated as a nested object, it has to be annotated with either `@o.a.s.c.solrj.beans.Field(child=true)` or `@o.s.d.s.core.mapping.ChildDocument`. The following uses the `@ChildDocument` annotation:

```
public class Book {  
  
    @Id String id;  
    @Indexed("type_s") String type;  
    @Indexed("title_t") String title;  
    @Indexed("author_s") String author;  
    @Indexed("publisher_s") String publisher;  
  
    @ChildDocument List<Review> reviews;      1  
  
    // setters and getters ...  
}  
  
public class Review {  
  
    @Id String id;                          2  
    @Indexed("type_s") String type;  
    @Indexed("review_dt") Date date;  
    @Indexed("stars_i") int stars;  
    @Indexed("author_s") String author;  
    @Indexed("comment_t") String comment;  
}
```

- 1 Multiple child documents can be associated with a parent one or use the domain type to store a single relationship.
- 2 Note that the nested document also needs to have a unique `id` assigned.

Assuming `Book#type` is `book`, and `Review#type` resolves to `review`, retrieving `Book` with its child relations `reviews` can be done by altering the `f1` query parameter, as the

following example shows:

```
Query query = new SimpleQuery(where("id").is("theWayOfKings"));
query.addProjectionOnField(new SimpleField("*"));
query.addProjectionOnField(new SimpleField("[child parentFilter=type_s:book]")); 1

return solrTemplate.queryForObject("books", query, Book.class);
```

- 1 The parent filter always defines the complete set of parent documents in the index, not the one for a single document.

Appendix

Appendix A: Namespace reference

The <repositories /> Element

The <repositories /> element triggers the setup of the Spring Data repository infrastructure. The most important attribute is `base-package`, which defines the package to scan for Spring Data repository interfaces. See [“XML configuration”](#). The following table describes the attributes of the <repositories /> element:

Table 4. Attributes

Name	Description
base-package	Defines the package to be scanned for repository interfaces that extend <code>*Repository</code> (the actual interface is determined by the specific Spring Data module) in auto-detection mode. All packages below the configured package are scanned, too. Wildcards are allowed.
repository-impl-postfix	Defines the postfix to autodetect custom repository implementations. Classes whose names end with the configured postfix are considered as candidates. Defaults to <code>Impl</code> .
query-lookup-strategy	Determines the strategy to be used to create finder queries. See “Query Lookup Strategies” for details. Defaults to <code>create-if-not-found</code> .

Name	Description
named-queries-location	Defines the location to search for a Properties file containing externally defined queries.
consider-nested-repositories	Whether nested repository interface definitions should be considered. Defaults to false.

Appendix B: Populators namespace reference

The <populator /> element

The <populator /> element allows to populate the a data store via the Spring Data repository infrastructure.^[1]

Table 5. Attributes

Name	Description
locations	Where to find the files to read the objects from the repository shall be populated with.

Appendix C: Repository query keywords

Supported query keywords

The following table lists the keywords generally supported by the Spring Data repository query derivation mechanism. However, consult the store-specific documentation for the exact list of supported keywords, because some keywords listed here might not be supported in a particular store.

Table 6. Query keywords

Logical keyword	Keyword expressions
AND	And
OR	Or
AFTER	After , IsAfter


Logical keyword	Keyword expressions
BEFORE	Before , IsBefore
CONTAINING	Containing, IsContaining, Contains
BETWEEN	Between , IsBetween
ENDING_WITH	EndingWith, IsEndingWith, EndsWith
EXISTS	Exists
FALSE	False , IsFalse
GREATER_THAN	GreaterThan , IsGreaterThan
GREATER_THAN_EQUALS	GreaterThanEqual , IsGreaterThanEqual
IN	In , IsIn
IS	Is , Equals , (or no keyword)
IS_EMPTY	IsEmpty , Empty
IS_NOT_EMPTY	IsNotEmpty , NotEmpty
IS_NOT_NULL	NotNull , IsNotNull
IS_NULL	Null , IsNull
LESS_THAN	LessThan , IsLessThan
LESS_THAN_EQUAL	LessThanEqual , IsLessThanEqual
LIKE	Like , IsLike
NEAR	Near , IsNear
NOT	Not , IsNot
NOT_IN	NotIn , IsNotIn
NOT_LIKE	NotLike , IsNotLike
REGEX	Regex , MatchesRegex , Matches

Logical keyword	Keyword expressions
STARTING_WITH	StartingWith , IsStartingWith , StartsWith
TRUE	True , IsTrue
WITHIN	Within , IsWithin

Appendix D: Repository query return types

Supported Query Return Types

The following table lists the return types generally supported by Spring Data repositories. However, consult the store-specific documentation for the exact list of supported return types, because some types listed here might not be supported in a particular store.



Geospatial types (such as `GeoResult` , `GeoResults` , and `GeoPage`) are available only for data stores that support geospatial queries.

Table 7. Query return types

Return type	Description
void	Denotes no return value.
Primitives	Java primitives.
Wrapper types	Java wrapper types.
T	An unique entity. Expects the query method to return one result at most. If no result is found, <code>null</code> is returned. More than one result triggers an <code>IncorrectResultSizeDataAccessException</code> .
Iterator<T>	An Iterator.
Collection<T>	A Collection.

Return type	Description
List<T>	A List.
Optional<T>	A Java 8 or Guava <code>Optional</code> . Expects the query method to return one result at most. If no result is found, <code>Optional.empty()</code> or <code>Optional.absent()</code> is returned. More than one result triggers an <code>IncorrectResultSizeDataAccessException</code> .
Option<T>	Either a Scala or Javaslang <code>Option</code> type. Semantically the same behavior as Java 8's <code>Optional</code> , described earlier.
Stream<T>	A Java 8 <code>Stream</code> .
Future<T>	A <code>Future</code> . Expects a method to be annotated with <code>@Async</code> and requires Spring's asynchronous method execution capability to be enabled.
CompletableFuture<T>	A Java 8 <code>CompletableFuture</code> . Expects a method to be annotated with <code>@Async</code> and requires Spring's asynchronous method execution capability to be enabled.
ListenableFuture	A <code>org.springframework.util.concurrent.ListenableFuture</code> . Expects a method to be annotated with <code>@Async</code> and requires Spring's asynchronous method execution capability to be enabled.
Slice	A sized chunk of data with an indication of whether there is more data available. Requires a <code>Pageable</code> method parameter.
Page<T>	A <code>Slice</code> with additional information, such as the total number of results. Requires a <code>Pageable</code> method parameter.
GeoResult<T>	A result entry with additional information, such as the distance to a reference location.
GeoResults<T>	A list of <code>GeoResult<T></code> with additional information, such as the average distance to a reference location.
GeoPage<T>	A <code>Page</code> with <code>GeoResult<T></code> , such as the average distance to a reference location.

Return type	Description
<code>Mono<T></code>	A Project Reactor <code>Mono</code> emitting zero or one element using reactive repositories. Expects the query method to return one result at most. If no result is found, <code>Mono.empty()</code> is returned. More than one result triggers an <code>IncorrectResultSizeDataAccessException</code> .
<code>Flux<T></code>	A Project Reactor <code>Flux</code> emitting zero, one, or many elements using reactive repositories. Queries returning <code>Flux</code> can emit also an infinite number of elements.
<code>Single<T></code>	A RxJava <code>Single</code> emitting a single element using reactive repositories. Expects the query method to return one result at most. If no result is found, <code>Mono.empty()</code> is returned. More than one result triggers an <code>IncorrectResultSizeDataAccessException</code> .
<code>Maybe<T></code>	A RxJava <code>Maybe</code> emitting zero or one element using reactive repositories. Expects the query method to return one result at most. If no result is found, <code>Mono.empty()</code> is returned. More than one result triggers an <code>IncorrectResultSizeDataAccessException</code> .
<code>Flowable<T></code>	A RxJava <code>Flowable</code> emitting zero, one, or many elements using reactive repositories. Queries returning <code>Flowable</code> can emit also an infinite number of elements.

1. see [XML configuration](#)

Version 4.0.8.RELEASE

Last updated 2019-05-13 15:01:58 MESZ